

Intel 8086



CSC 236

IBM PC body of knowledge 1978-present

Hardware

- 8086
- 80186
- 80286
- 80386
- 80486
- Pentium 1-4
- Core i3, i5, i7, i9

Software

- DOS
- Windows 1-3
- Windows 95, 98
- OS/2 Warp
- NT, 2000, XP
- Vista, 7
- 8, 10

**Why do I have
to understand
history?**



**Core i7
memory
hex word ■
0001**

**Java or C code:
short int x = ?**

What decimal value creates

hex value = 0001

Binary = 0000 0000 0000 0001

This is **not** a
trick question

**Core i7
memory
hex word =
0001**

**Java or C code:
short int x = 1**

What decimal value creates

hex value = 0001

Binary = 0000 0000 0000 0001

Wrong!

**Core i7
memory
hex word =
0001**

**Java or C code:
short int x = ?**

What decimal value creates

hex value = 0001

Binary = 0000 0000 000

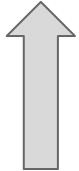
To know the correct
answer we need to
understand history

Decisions that shaped the PC revolution


IBM	Motorola	Intel



Decisions that shaped the PC revolution

IBM	Motorola	Intel
Internal uprocessor		
 Largest manufacturer of hardware and software		

Decisions that shaped the PC revolution

IBM	Motorola	Intel
Internal uprocessor		
Purchase <ul style="list-style-type: none">• hardware →		

Decisions that shaped the PC revolution

IBM	Motorola	Intel
Internal uprocessor		
Purchase <ul style="list-style-type: none">• hardware• software		



Decisions that shaped the PC revolution

IBM	Motorola	Intel
Internal uprocessor		
Purchase <ul style="list-style-type: none">• hardware• software	Why would the largest producer of	

chips and software
BUY
chips and software

Decisions that shaped the PC revolution

IBM	Motorola	Intel
Internal uprocessor		
Purchase <ul style="list-style-type: none">• hardware• software		

IBM sold \$10M mainframes

Target price for PCs ~\$1K

Decisions that shaped the PC revolution

IBM	Motorola	Intel
Internal uprocessor	16-bit 6800	
Purchase <ul style="list-style-type: none">• hardware• software	New start <ul style="list-style-type: none">• 68000• Full 32-bit• Excellent architecture• ... but developers must re-write their software	

Decisions that shaped the PC revolution

IBM	Motorola	Intel
Internal uprocessor	16-bit 6800	16-bit 8080
Purchase <ul style="list-style-type: none">● hardware● software	New start <ul style="list-style-type: none">● 68000● Full 32-bit● Excellent architecture● ... but developers must re-write their software	Conflicting requirements <ul style="list-style-type: none">● Preserve<ul style="list-style-type: none">○ 8080 code○ 16-bit regs● support > 64KB memory



Conundrum




(a confusing and difficult problem)

Conundrum

How can a program written for 64K memory and 16 bit registers use a whole Megabyte of memory?





Segmentation



**Divide memory into blocks, called
segments**

(16 to 64K bytes each)

A program can reference

- **CS — code segment**
- **DS — data segment**
- **SS — stack segment**

**3 segments \Rightarrow 2
bits**

**Why not have 4
segments?**

Segmentation



**Divide memory into blocks, called
segments**

(16 to 64K bytes each)

A program can reference

- **CS** — code segment
- **DS** — data segment
- **SS** — stack segment
- **ES** — extra segment



How segmentation works

1MB ranges from

- 00000_{16} **to**
- $FFFFF_{16}$

**Segment address is
5 hex digits**

**But 16-bits just give
us 4 digits**



How segmentation works


1MB ranges from

- 00000_{16} **to**
- $FFFF_{16}$

**Segment address is 5
hex digits**

**But 16-bits just give
us 4 digits**

Force lowest to be 0



How segmentation works


1MB ranges from

- 00000_{16} **to**
- $FFFFF_{16}$

Segment address is 5 hex digits

But 16-bits just give us 4 digits

Force lowest to be 0




**16-bit segments
register holds 4 hex
digits**

x x x x 0

**processor
implicitly adds
0**

How segmentation works

- CPU uses a segment register
 - 1 of 4 hardware registers
- On every memory reference
 - Instructions and data
 - Hardware calculates the absolute address
- Absolute address
 - The actual location in 1MB memory space
 - The start-of-segment plus a relative address
 - That's the real address



Summary

- Low-order hex digit of 5-hex digit (20-bit) segment address is always 0
- The smallest possible segment is 16 bytes
- Hardware calculates address for every reference
 - Instructions and data
 - $\text{ADD } x, y, z \leftarrow 4$ memory references
- Data cannot exceed 64KB
 - Without changing segment register value

Dynamic memory

- How can we be sure program does not violate 64KB limit?
- Static memory
 - Assign data to a segment
 - Use appropriate segment register
 - Can add many instructions
- Dynamic memory allocation
 - Requires runtime checks
 - Slows the program down by at least 10x

Who manages DS register?

- Programmer
 - Or compiler

Who manages DS register?

- Programmer
 - Or compiler
- Ruh Roh

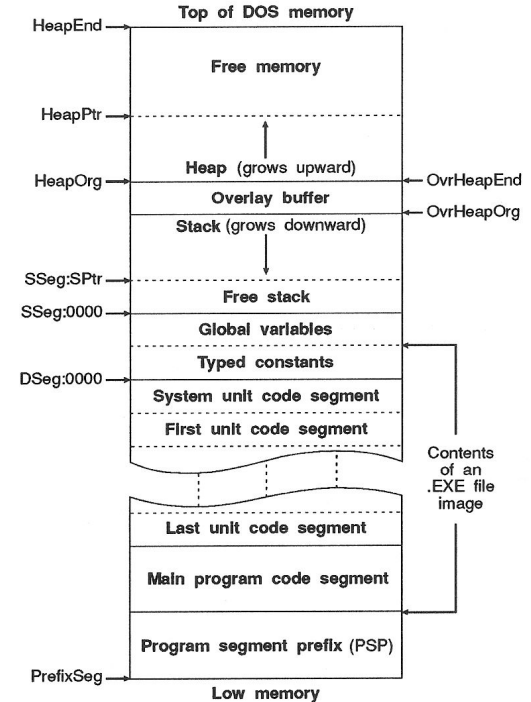


Who manages DS regis

- Programmer
 - Or compiler
- Ruh Roh

Early compiler documentation

Figure 16.1
Turbo Pascal memory map



The DS register is never changed during program execution. The size of the data segment cannot exceed 64K.

The data segment (addressed through DS) contains all typed constants followed by all global variables. The DS register is never changed during program execution. The size of the data segment cannot exceed 64K.


On entry to the program, the stack segment register (SS) and the stack pointer (SP) are loaded so that SS:SP points to the first byte past the stack segment. The SS register is never changed during

In 1985

- Intel released the 80386
- True 32-bit processor
- Max segment size 4GB
- Removed the need to update the segment registers



Is there anything else to worry about?



Begin Tangent



Gulliver's Travels

All Lilliputians must break their eggs only on the little end.

Those who broke eggs at the big end were angry.

Civil war broke out between

- the Little-Endians and
- the Big-Endians.



GULLIVER IN LILLIPUT.

End Tangent


How data is stored in memory

- Instruction

ADD 1000

A5	03	E8
----	----	----

- $1000_{10} = 03E8_{16}$



A problem




A problem

- How to add two 16-bit values

Add 1FA2 from memory

To 1234 in the CPU

XXXX




**Can only move 1
byte at a time**

A problem

- How to add two 16-bit values

Add 1FA2 from memory
To 1234 in the CPU
XXXX



**When adding 2-byte values,
want to move least
significant byte**


A problem

- How to add two 16-bit values

Add 1FA2 from memory

To 1234 in the CPU

XXXX



**When adding 2-byte values,
want to move least
significant byte.**

Must skip first byte


A problem

- How to add two 16-bit values

Add 1FA2 from memory

To 1234 in the CPU

XXXX



**When adding 2-byte values,
want to move least
significant byte.**

Must skip first byte




A problem

- How to add two 16-bit values

Add 1FA2 from memory

To 1234 in the CPU

XXXX



Then go back to first byte.

One-byte data bus

- Limitation of the 8080
 - Remember 8086 was an upgrade to 8080
- Data in memory
 - Store lowest-order byte first
 - “Little Endian”
 - To avoid skipping memory
- In order to maintain backward compatibility
 - 8086 (which is 16-bit) is little endian

Little Endian

ADD 1000

A5	E8	03
----	----	----

03
E8
A5

Byte swapping

Big Endian

ADD 1000

A5	03	E8
----	----	----

E8
03
A5

$$1000_{10} = 03E8_{16}$$

Problems with Little Endian

- Sharing data with Big Endian machines
- When humans read code

Hex 00 01 in memory is 256_{10} in CPU

00 01 swapped is 01 00 = 256_{10}

Problems with Little Endian

- Sharing data with Big Endian machines
- When humans read code
- Latest CPUs (from Intel) are still little endian
 - Even though it has been 40 years since last byte-wide memory chip

Byte reversal only occurs in memory

- Declare a word variable of 1000_{10}
- Created value is $03E8_{16}$
- Stored in memory as $E803_{16}$
- Reversed whenever moved between memory and CPU
 - Either direction
 - Always

Byte reversal only occurs in memory

- Only swap multi-byte data
- Byte: 0A \Rightarrow 0A
- Word: 0102 \Rightarrow 0201
- Double word: 01020304 \Rightarrow 04030201
- String: “abc”
 - 616263 (in ASCII hex) becomes
 - 616263
 - No swapping because each character is one byte wide
- Swapping is context sensitive

8086 Programmer's Model

8086 Programmer's Model

- How the compiler writer or assembler programmer views the system
- Starts with the register set

8086 Programmer's Model

- How the compiler writer or assembler programmer views the system
- Starts with the register set
 - Special purpose registers

Instruction ptr	IP
-----------------	-----------

Code segment	CS
Data segment	DS
Stack segment	SS
Extra segment	ES

Stack pointer	SP
Base pointer	BP
Source index	SI
Destination idx	DI

8086 Programmer's Model

- How the compiler writer or assembler programmer views the system
- Starts with the register set
 - Special purpose registers

Instruction ptr

IP

Code segment

CS

Data segment

DS

Stack segment

SS

Extra segment

ES

Stack pointer

SP

Base pointer

BP

Source index

SI

Destination idx

DI

8086 Programmer's Model

- How the compiler writer or assembler programmer views the system
- Starts with the register set
 - Special purpose registers

Code segment	CS
Data segment	DS
Stack segment	SS
Extra segment	ES

Instruction ptr	IP
-----------------	-----------

Stack pointer	SP
Base pointer	BP
Source index	SI
Destination idx	DI

8086 Programmer's Model

- How the compiler writer or assembler programmer views the system
- Starts with the register set
 - Special purpose registers

Code segment	CS
Data segment	DS
Stack segment	SS
Extra segment	ES

Instruction ptr	IP
-----------------	-----------

Stack pointer	SP
Base pointer	BP
Source index	SI
Destination idx	DI

8086 Programmer's Model

- How the compiler writer or assembler programmer views the system
- Starts with the register set
 - Special purpose registers

Code segment
Data segment
Stack segment
Extra segment

CS
DS
SS
ES

Instruction ptr

IP

Stack pointer
Base pointer
Source index
Destination idx

SP
BP
SI
DI

Segment registers

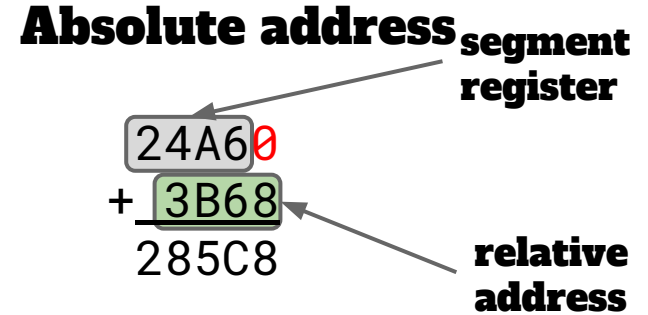
- Used to calculate absolute address
- Who sets segment registers
 - OS
 - programmer/compiler
- OS
 - CS } **program execution**
 - SS }
- Writer
 - DS } **data access**
 - ES }

Absolute address

$$\begin{array}{r} 24A60 \\ + 3B68 \\ \hline 285C8 \end{array}$$

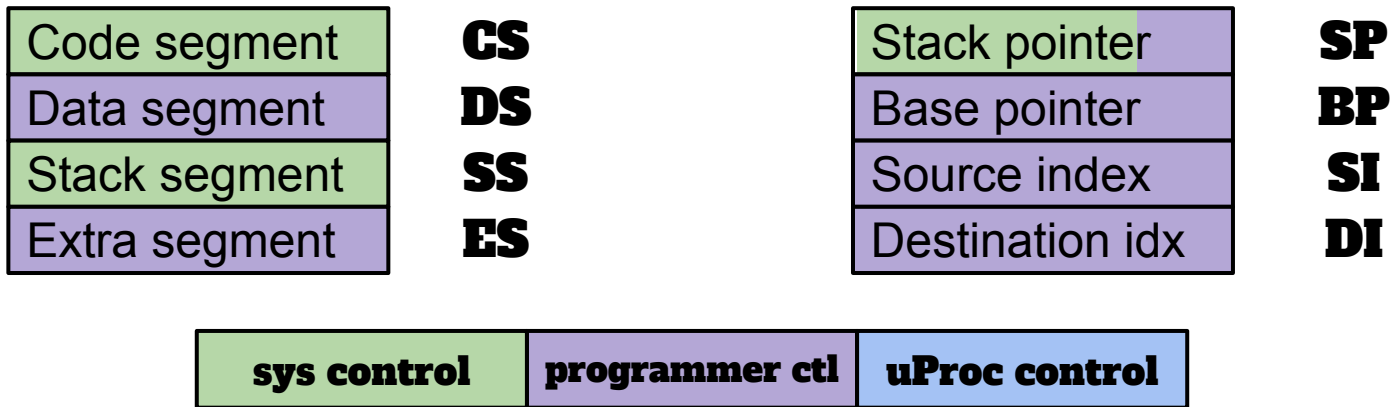
segment register

relative address



8086 Programmer's Model

- How the compiler writer or assembler programmer views the system
- Starts with the register set
 - Special purpose registers



8086 Programmer's Model

- General purpose registers
 - For use by the programmer
 - Byte and word sizes

16-bit registers

ax	accumulator
bx	base
cx	count
dx	data

8086 Programmer's Model

- General purpose registers
 - For use by the programmer
 - Byte and word sizes

When used as
8-bit register,
high and low
are
disconnected

8-bit registers

ah	al
bh	bl
ch	cl
dh	dl

16-bit registers

ax **accumulator**
bx **base**
cx **count**
dx **data**

GP registers

AX = FFFF

0001 Add 1 to AX

AX = 0000 All 16 bits
affected

AX = FFFF

01 Add 1 to AL

AX = FF00 Low 8 bits
affected

Condition codes

- Gives status of result
 - Negative
 - Overflow
 - Zero
 - ...

Condition codes

- SF — sign flag
 - SF = 0 — result of last operation is positive (or zero)
 - SF = 1 — result of last operation is negative
- ZF — zero flag
 - ZF = 1 — result of last operation is zero
 - ZF = 0 — result of last operation is not zero
- OF — overflow flag
 - OF = 1 — a signed overflow occurred
carry out of sign bit \neq carry in of sign bit
- CF — carry flag
 - CF = 1 — an unsigned overflow occurred
carry out (on add) or borrow (on subtract)

Significant question

- Intel add instruction works on
 - Unsigned and
 - Signed (2-comp)
- And hardware does not
 - Know or
 - Care
- How can it set the overflow and carry flags?
 - Answer: it answers both questions
 - The programmer has to check the relevant flags ... based on what they are trying to do.

Condition code summary

- SF — equals left most bit of result
- ZF — equals 1 iff all bits are zero
- OF — equals (carry out \neq carry in) of left-most bit
- CF — equals
 - Carry out of left-most bit on add
 - Borrow of left-most bit on sub

Use OF w/
signed values

Use CF w/
unsigned values

Example 1

	1111
AX	FFFF
Add to AX	<u>0001</u>
	0000

SF = 0

ZF = 1

OF = 0

CF = 1

Example 2

	11
AX	FF FF
Add to AL	<u>00 01</u>
	FF 00

SF = 0

ZF = 1

OF = 0

CF = 1

Example 3

AX	7000
Add to AX	<u>1000</u>
	8000

SF = 1

ZF = 0

OF = 1

CF = 0

Example 3

AX	7000
Add to AX	<u>1000</u>
	8000

111				
0111	0000	0000	0000	
<u>0001</u>	<u>0000</u>	<u>0000</u>	<u>0000</u>	
1000	0000	0000	0000	

SF = 1

ZF = 0

OF = 1

CF = 0

Example 4

	1111
AX	FFFF
Add to AX	<u>FFFF</u>
	FFFE

SF = 1

ZF = 0

OF = 0

CF = 1