

ARM



CSC 236



- Advanced RISC Machine
 - Family of microprocessors
 - Designed by ARM holdings
 - Licenses semiconductor design
 - ... but doesn't manufacture any chips
- Advantages
 - Low transistor count
 - Low power usage
 - Ideal for embedded systems

The ARM logo is displayed in the bottom right corner in a large, blue, lowercase, sans-serif font.

Dominates the embedded market

- 95% of phones
- 35% of hi-tech consumer devices
- Apple is using 8- to 10-core 5nm ARM processor in newest systems

Apple announced the M2 SoC on June 6, 2022, at [WWDC](#), along with the new MacBook Air and the new 13-inch MacBook Pro and later the [iPad Pro \(6th generation\)](#). It is the successor to the Apple M1. The M2 is made with TSMC's "Enhanced 5-nanometer technology" N5P process and contains 20 billion transistors, a 25% increase from the previous generation M1. The M2 can be configured with up to 24 gigabytes of RAM and 2 terabytes of storage. It has 8 CPU cores (4 performance and 4 efficiency) and up to 10 GPU cores. The M2 also increases the memory bandwidth to 100 GB/s. Apple claims CPU improvements up to 18% and GPU improvements up to 35% compared to the previous M1

Characteristics

- Load/store architecture
 - Data loaded from memory into registers
 - Operations on registers
 - Data stored from registers into memory
- Instructions
 - All are same size
 - Execute in one cycle^{*}
- Two powerful extensions compared to 8086
 - Conditional execution
 - Optionally set condition codes
- Many variants
 - We'll look at the 32-bit architecture

**(*) per pipeline stage
excluding load and store**








Programmer's model

- 4GB address space (that's 32 bits of addresses)
- Data size
 - Byte (8 bits)
 - Halfword (16 bit)
 - Word (32 bits)
- Registers
 - 16 registers, each 32-bit
 - General purpose: r0 - r12
 - SP: r13 — stack pointer
 - LR: r14 — link register (like BP)
 - PC: r15 — program counter
 - A status register for condition codes

Memory operations

- Load and store
 - Other instructions can't access memory
 - Source is unmodified.
- Different instructions for:
 - Data size
 - Signed/unsigned

Load

Instruction	Operands	Notes
LDR	rd, =var	rd  address of variable
LDR	rd, =const	rd  immediate constant
LDR	rd, [rn]	rd  value at rn
LDR	rd, [rn, rm]	rd  value at rn + rm
LDR	rd, [rn, #n]	rd  value at rn + #n
LDR	rd, [rn], #n	rd  value at rn; rn  rn + #n

No LDR r1, [var]

Instead:
LDR r1, =var
LDR r1, [r1]

I'm like a
post-increment.

Store

Notice, source and destination are reversed.

Instruction	Operands	Notes
STR	rs, [rn]	[rn] ← rs
STR	rs, [rn, rm]	[rn + rm] ← rs
STR	rs, [rn, #n]	[rn + #n] ← rs
STR	rs, [rn], #n	[rn] ← rs; rn ← rn + #n

Data Size

Instruction	Data size
LDR	Word size (32 bits)
LDRH	Halfword, load into low-order 16 bits, zero-fill high-order 16 bits
LDRB	Byte, load into low-order 8 bits, zero-fill high-order 24 bits
LDRSH	Signed halfword, load into low-order 16 bits and sign-extend to 32 bits
LDRSB	Signed byte, load into low-order 8 bits and sign-extend to 32 bits
STR	Word size (32 bits)
STRH	Halfword, from the low-order 16 bits.
STRB	Byte, from the low-order 8 bits.





Don't need signed versions of these.

Move

- Load and store are for memory operations
- Not for copying between registers.

Instruction n	Operands	Notes
MOV	rd, rn	rd ← rn

Add and subtract

Instruction	Operands	Notes
ADD	rd, rn, rm	rd  rn + rm
ADD	rd, rn, #n	rd  rn + #n
SUB	rd, rn, rm	rd  rn - rm
SUB	rd, rn, #n	rd  rn - #n

- Optionally set condition code

```
add r3,r2,r1 ;r3 = r2+r1  
           ;CC not set
```





```
adds r3,r2,r1 ;r3 = r2+r1  
          ;CC is set
```



suffix “s” says “set condition codes”

This applies to many instructions

Multiply and Divide





Instruction	Operands	Notes
UMULL	rdhi, rdlo, rm, rs	rdhi:rdlo  rm * rs
SMULL	rdhi, rdlo, rm, rs	rdhi:rdlo  rm * rs
UDIV	rd, rm, rn	rd  rn / rm
SDIV	rd, rm, rn	rd  rn / rm

Signed and Unsigned versions of these operations.

Multiply yields 64 bits

Divide yields 32 bits.

Multiply and Divide

Instruction	Operands	Notes
UMULL	rdhi, rdlo, rm, rs	rdhi:rdlo  rm * rs
SMULL	rdhi, rdlo, rm, rs	rdhi:rdlo  rm * rs
UDIV	rd, rm, rn	rd  rn / rm
SDIV	rd, rm, rn	rd  rn / rm

Some restrictions on what registers you can use here.

Divide may not be implemented in hardware.

Compare

Instruction	Operands	Notes
CMP	rd, rn	rd ? rn
CMP	rd, #n	rd ? #n

- Like x86
 - Performs subtraction
 - Discards result
- Always sets the condition codes

Status register

- 4 primary bits
 - Similar to Intel
 - N, Z, C, V — like SF, ZF, CF, OF
- N — negative
 - 0 \Rightarrow result is positive (or zero)
 - 1 \Rightarrow result is negative
 - Same as Intel SF
- Z — zero
 - 0 \Rightarrow result is not zero
 - 1 \Rightarrow result is zero
 - Same as Intel ZF
- V — signed overflow
 - 1 \Rightarrow signed overflow
 - 0 \Rightarrow no signed overflow
 - Same as Intel OF
- C — unsigned overflow
 - C = carry out
 - **Not exactly the same** as Intel CF

Direct subtraction implementation

0001
- 0010

Direct subtraction implementation

$$\begin{array}{r} 0001 \\ - 0010 \\ \hline 1 \end{array}$$

Direct subtraction implementation

$$\begin{array}{r} 0001 \\ - 0010 \\ \hline 1 \end{array}$$

Direct subtraction implementation

Assume leading

$$\begin{array}{r} \text{1} \\ \downarrow \\ \begin{array}{r} 10001 \\ - 0010 \\ \hline 1 \end{array} \end{array}$$

Direct subtraction implementation

Redistribute

$$\begin{array}{r} 1 \\ 0111 \\ +0001 \\ - \underline{0010} \\ 1 \end{array}$$

Direct subtraction implementation

Redistribute

```
    1
0111
+0001
- 0010
-----
1111
```

SF = 1

ZF = 0

OF =

CF =

Direct subtraction implementation

```
    1
0111
+0001
- 0010
-----
1111
```

SF = 1

ZF = 0

OF = +1 - 2 — no signed overflow

CF =

Direct subtraction implementation

1
0111
~~0001~~
- 0010
1111

SF = 1
ZF = 0
OF = 0
CF =



Direct subtraction implementation

$$\begin{array}{r} 1 \\ 0111 \\ +0001 \\ - 0010 \\ \hline 1111 \end{array}$$

SF = 1
ZF = 0
OF = 0
CF = 1

Did you need to
borrow into the
high-order bit.

Makes sense. 1 – 2 is
an unsigned overflow.

Status register

- 4 primary bits
 - Similar to Intel
 - N, Z, C, V — like SF, ZF, CF, OF
- N — negative
 - 0 \Rightarrow result is non-negative
 - 1 \Rightarrow result is negative
 - Same as Intel SF
- Z — zero
 - 0 \Rightarrow result is not zero
 - 1 \Rightarrow result is zero
 - Same as Intel ZF
- V — signed overflow
 - 1 \Rightarrow signed overflow
 - 0 \Rightarrow no signed overflow
 - Same as Intel OF
- C — unsigned overflow
 - C = carry out
 - **Not the same** as Intel CF
 - Semantics: depends on operation

Unsigned overflow	Intel	ARM
Add	CF = 1	C = 1
Sub	CF = 1	C = 0

Subtraction on ARM

$$\begin{array}{r} 1 \\ 0111 \\ +0001 \\ - 0010 \\ \hline 1111 \end{array}$$

N = 1
Z = 0
V = 0
C = 0

What was left in that extra bit after subtraction?

$$\begin{array}{r} 1 \\ 0001 \\ 10010 \\ - 0001 \\ \hline 0001 \end{array}$$

N = 0
Z = 0
V = 0
C = 1

OK. It was still 1 after subtraction.

		Flags	Arithmetic	Compare
SIGNED				
PL	Plus	N=0	Positive result	
MI	Minus	N=1	Negative result	
VC	Overflow clear	V=0	No signed ovfl	
VS	Overflow set	V=1	signed overflow	
GE	Greater / equal	N=V	result ≥ 0	dest \geq source
GT	Greater	Z=0 & N=V	result > 0	dest $>$ source
LE	Less / equal	Z=1 or N \neq V	result ≤ 0	dest \leq source
LT	Less	N \neq V	result < 0	dest $<$ source

dest - source

These are the names for these comparisons.

		Flags	Arithmetic	Compare
SIGNED				
PL	Plus	N=0	Positive result	
MI	Minus	N=1	Negative result	
VC	Overflow clear	V=0	No signed ovfl	
VS	Overflow set	V=1	signed overflow	
GE	Greater / equal	N=V	result ≥ 0	dest \geq source
GT	Greater	Z=0 & N=V	result > 0	dest $>$ source
LE	Less / equal	Z=1 or N \neq V	result ≤ 0	dest \leq source
LT	Less	N \neq V	result < 0	dest $<$ source

dest - source

Result should be non-negative .. unless there's overflow

Same as 8086

		Flags	Arithmetic	Compare
SIGNED				
PL	Plus	N=0	Positive result	
MI	Minus	N=1	Negative result	
VC	Overflow clear	V=0	No signed ovfl	
VS	Overflow set	V=1	signed overflow	
GE	Greater / equal	N=V	result ≥ 0	dest \geq source
GT	Greater	Z=0 & N=V	result > 0	dest $>$ source
LE	Less / equal	Z=1 or N \neq V	result ≤ 0	dest \leq source
LT	Less	N \neq V	result < 0	dest $<$ source

dest - source

Negation of GE
Condition

		Flags	Arithmetic	Compare
SIGNED				
PL	Plus	N=0	Positive result	
MI	Minus	N=1	Negative result	
VC	Overflow clear	V=0	No signed ovfl	
VS	Overflow set	V=1	signed overflow	
GE	Greater / equal	N=V	result ≥ 0	dest \geq source
GT	Greater	Z=0 & N=V	result > 0	dest $>$ source
LE	Less / equal	Z=1 or N \neq V	result ≤ 0	dest \leq source
LT	Less	N \neq V	result < 0	dest $<$ source

dest - source

GE Test, but excluding zero.

Negation of GT test.

		Flags	Arithmetic	Compare
SIGNED				
PL	Plus	N=0	Positive result	
MI	Minus	N=1	Negative result	
VC	Overflow clear	V=0	No signed ovfl	
VS	Overflow set	V=1	signed overflow	
GE	Greater / equal	N=V	result ≥ 0	dest \geq source
GT	Greater	Z=0 & N=V	result > 0	dest $>$ source
LE	Less / equal	Z=1 or N \neq V	result ≤ 0	dest \leq source
LT	Less	N \neq V	result < 0	dest $<$ source

Or, after an arithmetic operation

The condition codes tell you what you got.

		Flags	Arithmetic	Compare
UNSIGNED				
CS	Carry set	C=1	Add overflow	dest \geq source
CC	Carry clear	C=0	Sub overflow	dest < source
HI	High	C=1 & Z=0	Sub result > 0	dest > source
LS	Low or same	C=0 or Z=1	Sub result \leq 0	dest \leq source
BOTH SIGNED AND UNSIGNED				
EQ	Equal	Z=1	Equal	dest = source
NE	Not equal	Z=0	Not equal	dest \neq source
AL	Always		Unconditional	

dest - source

Remember, carry is
backward on
subtraction

		Flags	Arithmetic	Compare
UNSIGNED				
CS	Carry set	C=1	Add overflow	dest \geq source
CC	Carry clear	C=0	Sub overflow	dest < source
HI	High	C=1 & Z=0	Sub result > 0	dest > source
LS	Low or same	C=0 or Z=1	Sub result \leq 0	dest \leq source
BOTH SIGNED AND UNSIGNED				
EQ	Equal	Z=1	Equal	dest = source
NE	Not equal	Z=0	Not equal	dest \neq source
AL	Always		Unconditional	

dest - source

Negation of the CS
(\geq) test.

		Flags	Arithmetic	Compare
UNSIGNED				
CS	Carry set	C=1	Add overflow	dest \geq source
CC	Carry clear	C=0	Sub overflow	dest < source
HI	High	C=1 & Z=0	Sub result > 0	dest > source
LS	Low or same	C=0 or Z=1	Sub result \leq 0	dest \leq source
BOTH SIGNED AND UNSIGNED				
EQ	Equal	Z=1	Equal	dest = source
NE	Not equal	Z=0	Not equal	dest \neq source
AL	Always		Unconditional	

dest - source

CS test (\geq), excluding zero.

Negation of HI (>) test.

		Flags	Arithmetic	Compare
UNSIGNED				
CS	Carry set	C=1	Add overflow	dest \geq source
CC	Carry clear	C=0	Sub overflow	dest < source
HI	High	C=1 & Z=0	Sub result > 0	dest > source
LS	Low or same	C=0 or Z=1	Sub result \leq 0	dest \leq source
BOTH SIGNED AND UNSIGNED				
EQ	Equal	Z=1	Equal	dest = source
NE	Not equal	Z=0	Not equal	dest \neq source
AL	Always		Unconditional	

Or, after an unsigned arithmetic operation.

This tells you what you got.

		Flags	Arithmetic	Compare
UNSIGNED				
CS	Carry set	C=1	Add overflow	dest \geq source
CC	Carry clear	C=0	Sub overflow	dest < source
HI	High	C=1 & Z=0	Sub result > 0	dest > source
LS	Low or same	C=0 or Z=1	Sub result \leq 0	dest \leq source
BOTH SIGNED AND UNSIGNED				
EQ	Equal	Z=1	Equal	dest = source
NE	Not equal	Z=0	Not equal	dest \neq source
AL	Always		Unconditional	

These tests make sense for signed or unsigned.

Branch Instructions

Instruction	Operands	Notes
beq	label	Branch if equal (Z=1)
bne	label	Branch if not equal (Z=0)
bcs	label	Branch if carry set (C=1, overflow on add, no overflow on subtract, dest \geq source)
bpl	label	Branch if result is positive (N=0)
bge	label	Branch if greater than (Z=0 and N=V)
bal	label	Always jump

We can use labels,
just like with MASM.

A branch instruction
for every condition,
including “Always”

Conditional execution

- Goal: `abs(r4)`

Conditional execution

- Goal: `abs(r4)`
 - Depends on `r0 = 0`
 - `(ldr r0, #0)`

Conditional execution

- Goal: `abs(r4)`
 - Depends on `r0 = 0`
 - `(ldr r0, #0)`
- Code

```
cmp r4, #0
```


Conditional execution

- Goal: `abs(r4)`
 - Depends on `r0 = 0`
 - `(ldr r0, #0)`
- Code

```
cmp r4, #0  
bpl skip
```

Branch if plus (`result > 0`)

Conditional execution

- Goal: `abs(r4)`
 - Depends on `r0 = 0`
 - `(ldr r0, #0)`
- Code

```
cmp r4, #0
bpl skip
sub r4, r0, r4
```

```
skip:
```

Conditional execution

- Goal: abs(r4)
 - Depends on r0 = 0
 - (ldr r0, #0)
- Code

```
cmp r4, #0
bpl skip
sub r4, r0, r4
```

skip:

- Condition execution
 - Execute sub conditionally
 - If result < 0
- Code

```
cmp r4, #0
bpl skip
submi r4, r0, r4
```

Execute sub if minus (result < 0)

Conditional execution

- Goal: abs(r4)
 - Depends on $r0 = 0$
 - (ldr r0, #0)
- Code

```
cmp r4, #0
bpl skip
sub r4, r0, r4
```

skip:

- Condition execution
 - Execute sub conditionally
 - If result < 0
- Code

```
cmp r4, #0
bpl skip
submi r4, r0, r4
```

These are the same names used by the conditionals.

Conditional execution

- Fewer instructions
 - Yes! That's good.
 - But the real reason for this is ...
 - ... pause for architecture lesson

Pipelining

- Modern architectures are pipelined
 - Execution is divided into stages
 - Instructions pass from stage to stage
 - Like cars on an assembly line
 - Multiple instructions “in flight”
 - Increases parallelism
 - Working on multiple instructions at once
 - Increases clock speed
 - Less to do in each pipeline stage.



Pipelining example

- Suppose there are 3 stages

- Fetch
- Decode
- Execute

Early ARM had 3-stage pipeline

- Fetch

- Gets instruction and operands (if any) from memory

- Decode

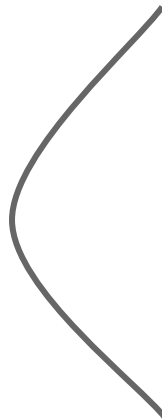
- Executes logic to understand and dispatch instruction

- Execute

- Performs requested operation

Pipelining example

	1	2	3	4	5	6	7	8
LDR								
ADD								
BEQ								
MOV								
STR								
CMP								



instructions

time

Pipelining example

	1	2	3	4	5	6	7	8
LDR	F							
ADD								
BEQ								
MOV								
STR								
CMP								

Pipelining example

	1	2	3	4	5	6	7	8
LDR	F	D						
ADD		F						
BEQ								
MOV								
STR								
CMP								



Pipelining example

	1	2	3	4	5	6	7	8
LDR	F	D	E					
ADD		F	D					
BEQ			F					
MOV								
STR								
CMP								



Pipelining example

	1	2	3	4	5	6	7	8
LDR	F	D	E					
ADD		F	D					
BEQ			F					
MOV								
STR								
CMP								

**What happens now?
Which instruction is
next?
MOV?
CMP?**

Pipelining example

	1	2	3	4	5	6	7	8
LDR	F	D	E					
ADD		F	D					
BEQ			F					
MOV								
STR								
CMP								

**ARM assumes branch
NOT taken;
fetches MOV
instruction**

Pipelining example

	1	2	3	4	5	6	7	8
LDR	F	D	E					
ADD		F	D	E				
BEQ			F	D				
MOV				F				
STR								
CMP								

Pipelining example

	1	2	3	4	5	6	7	8
LDR	F	D	E					
ADD		F	D	E				
BEQ			F	D	E			
MOV				F	D			
STR					F			
CMP								

Knows now if branch was taken

If not taken; keep going

What if branch is taken?

Pipelining example

	1	2	3	4	5	6	7	8
LDR	F	D	E					
ADD		F	D	E				
BEQ			F	D	E			
MOV				F	D			
STR					F			
CMP								

Do not want these

- **Flush pipeline**
- **Undo any effects**
- **No effects in this simple pipeline**

Pipelining example

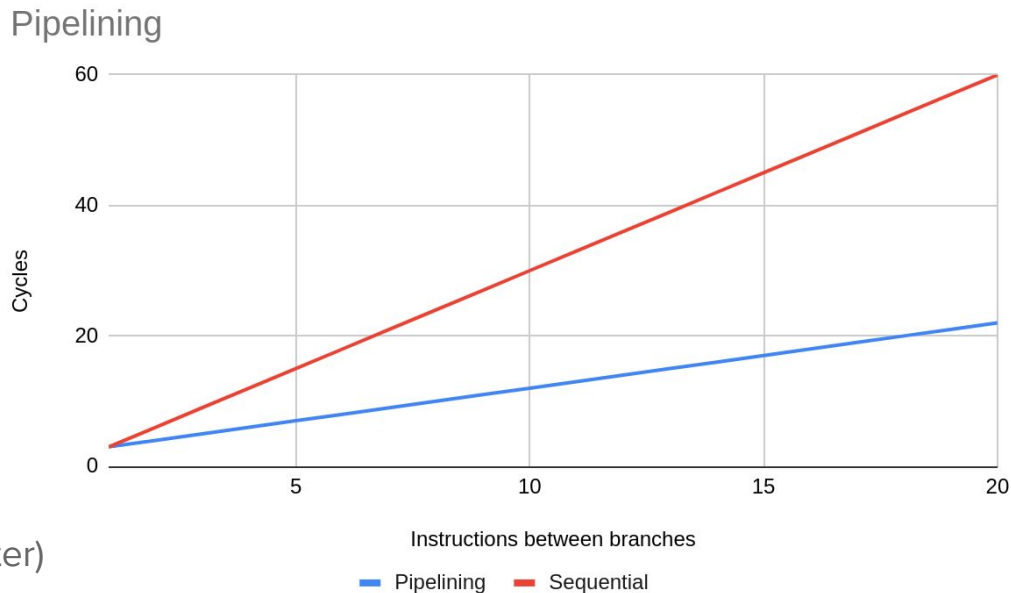
	1	2	3	4	5	6	7	8
LDR	F	D	E					
ADD		F	D	E				
BEQ			F	D	E			
MOV				F	D			
STR					F			
CMP						F		

Do not want these

- **Flush pipeline**
- **Undo any effects**
- **No effects in this simple pipeline**

Pipelining performance

- One instruction
 - 3 cycles to complete
- N instructions
 - Without flushing pipeline
 - $N + 2$ cycles to complete
 - Not $3N$
- Performance enhancement
 - Clock speed is ~ 3 times faster (shorter)
 - If average N instructions between branches (flushes)
 - $N + 2$ versus $3N$



Pipelining abs(r4)

	1	2	3	4	5	6	7	8
LDR								
CMP								
SUBMI								
MOV								
STR								
...								
...								
...								

Pipelining abs(r4)

	1	2	3	4	5	6	7	8
LDR	F	D	E					
CMP		F	D					
SUBMI			F					
MOV								
STR								
...								
...								
...								

Pipelining abs(r4)

	1	2	3	4	5	6	7	8
LDR	F	D	E					
CMP		F	D	E				
SUBMI			F	D	E			
MOV				F	D			
STR					F			
...								
...								
...								

Only execute if mi (minus)

Conditional Execution

	Conditionally executed	Optionally set CC
ADD / SUB	Yes	Yes — N,Z,C,V
AND / ORR / EOR	Yes	Yes — N,Z,C
BXX	Yes	No
CMP	Yes	Always
LDR / STR	Yes	No
MOV	Yes	Yes — N,Z,C
STMDB / LDMIA	Yes	No
UMUL / SMULL	Yes	Yes — N,Z

- How to code
 - Conditional execution
 - Set CC

- Format

opcode**cond****set**

- **cond** — Execution depends on CC

- `addmi`
- `subpl`

- **set** — 's'

- `adds`
- `subs`

Can be combined:
`addmis`

Immediate data

- add r0, r0, #1 r0=r0 +1
- add r0, r0, #10 r0=r0 +10
- add r0, r0, #100 r0=r0 +100
- add r0, r0, #999 illegal
- add r0, r0, #1000 r0=r0 +1000
- add r0, r0, #1001 illegal

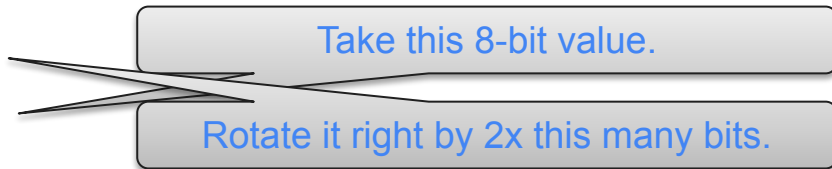
Immediate data

- X86 operands can be encoded in additional bytes.
 - If immediate is byte — data component is 1 byte
 - If immediate is word — data component is 2 bytes
 - Variable length instructions
- ARM machine instructions are always 32 bits
- Bits are needed for
 - Opcode
 - Operands
- We can't spend all 32 bits on an immediate value ... what to do?

Immediate data

- ARM uses 12 bits for immediate value

- Rotate — 4 bits 0-15
- Value — 8 bits 0-255



- Why rotate **right**?

- Rotate left by N bits is just rotate right by $32 - N$.

- This produces a 32-bit value

- But just some 32-bit values.
- ... hopefully, the ones you need most often.

Immediate data

ROT	F	E	D	C	...	1	0
2 x ROT	30	28	26	24		2	0
Rot left	2	4	6	8		30	0
Multiply	4	16	64	256		2^{30}	1

Immediate data

ROT	F	E	D	C	...	1	0
2 x ROT	30	28	26	24		2	0
Rot left	2	4	6	8		30	0
Multiply	4	16	64	256		2^{30}	1

0000 0000 0000 0000 0000 0000 abcd efgh

Any sequence
of 8 bits.

Immediate data

ROT	F	E	D	C	...	1	0
2 x ROT	30	28	26	24		2	0
Rot left	2	4	6	8		30	0
Multiply	4	16	64	256		2^{30}	1

0000 0000 0000 0000 0000 0000 abcd efgh

0000 0000 0000 0000 0000 abcd efgh 0000

Immediate data

ROT	F	E	D	C	...	1	0
2 x ROT	30	28	26	24		2	0
Rot left	2	4	6	8		30	0
Multiply	4	16	64	256		2^{30}	1

0000 0000 0000 0000 0000 0000 abcd efgh

0000 0000 0000 0000 abcd efgh 0000 0000

Immediate data

ROT	F	E	D	C	...	1	0
2 x ROT	30	28	26	24		2	0
Rot left	2	4	6	8		30	0
Multiply	4	16	64	256		2^{30}	1

0000 0000 0000 0000 0000 0000 abcd efgh

gh00 0000 0000 0000 0000 0000 00ab cdef

Immediate data

ROT	F	E	D	C	...	1	0
2 x ROT	30	28	26	24		2	0
Rot left	2	4	6	8		30	0
Multiply	4	16	64	256		2^{30}	1

- **Generate #1000**
 - **$1000 = 111101000_2 = 1111010_2 * 4$**
 - **Rotate = F**
 - **Value = $1111010_2 = FA$**

Immediate data

ROT	F	E	D	C	...	1	0
2 x ROT	30	28	26	24		2	0
Rot left	2	4	6	8		30	0
Multiply	4	16	64	256		2^{30}	1

- **Cannot generate #999**
- **Cannot generate #1001**

(that's 11110011₂)
(that's 111101001₂)

