

# JVM




CSC 236

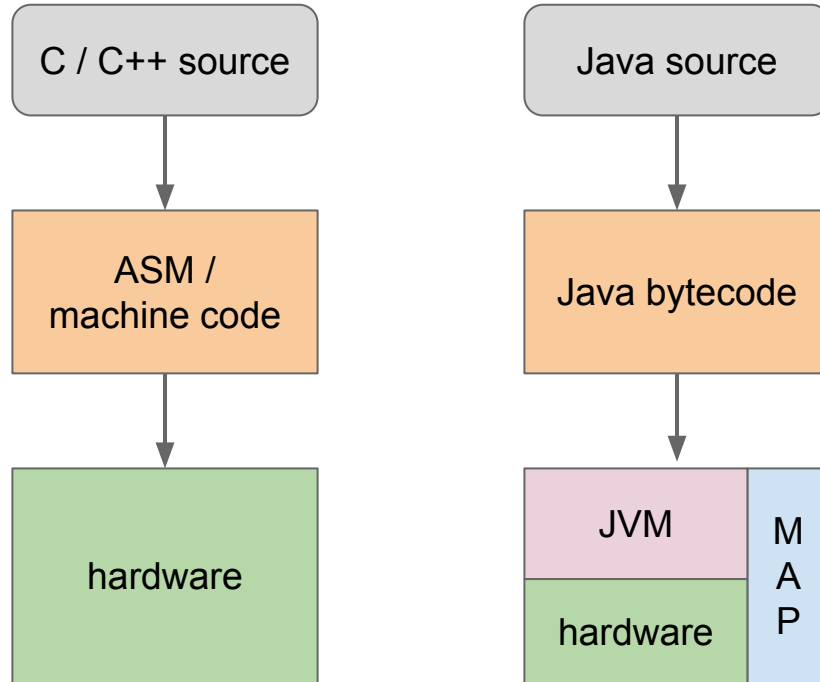
# Java virtual machine

- Java compile
  - Generates *bytecode*
  - From source
- JVM execute
  - Bytecode
  - On any machine
- Motto: write once run anywhere
  - Eliminate problem of incompatible executables

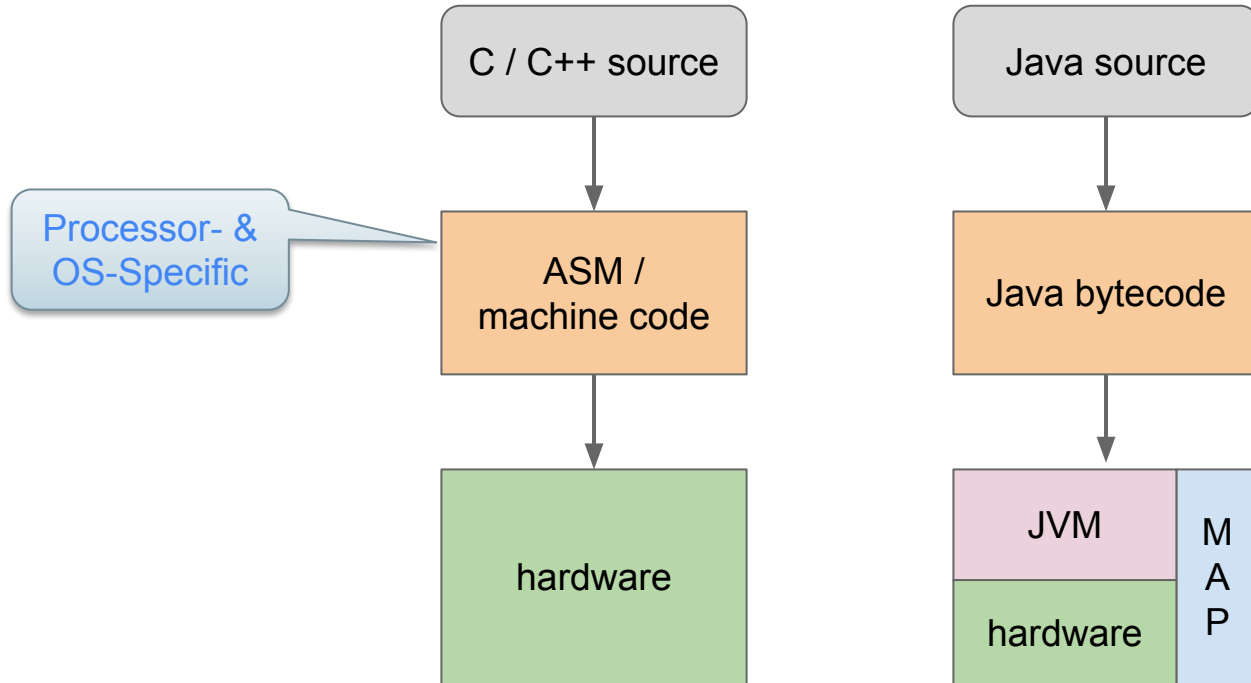
# JVM — Java Virtual Machine




# JVM — Java Virtual Machine



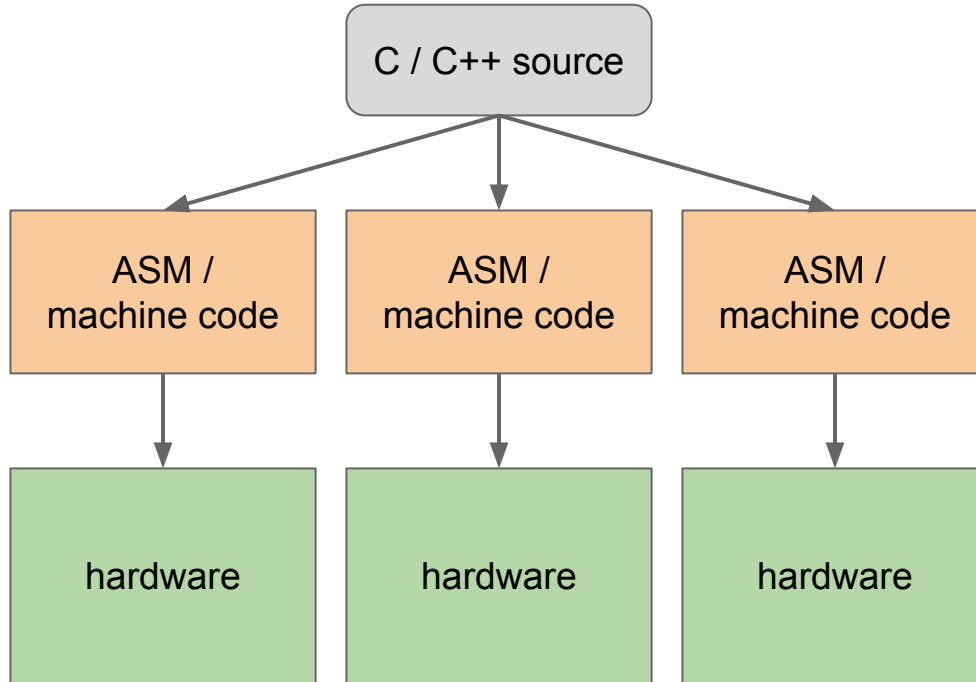
# Traditional (compiled) language vs. Java



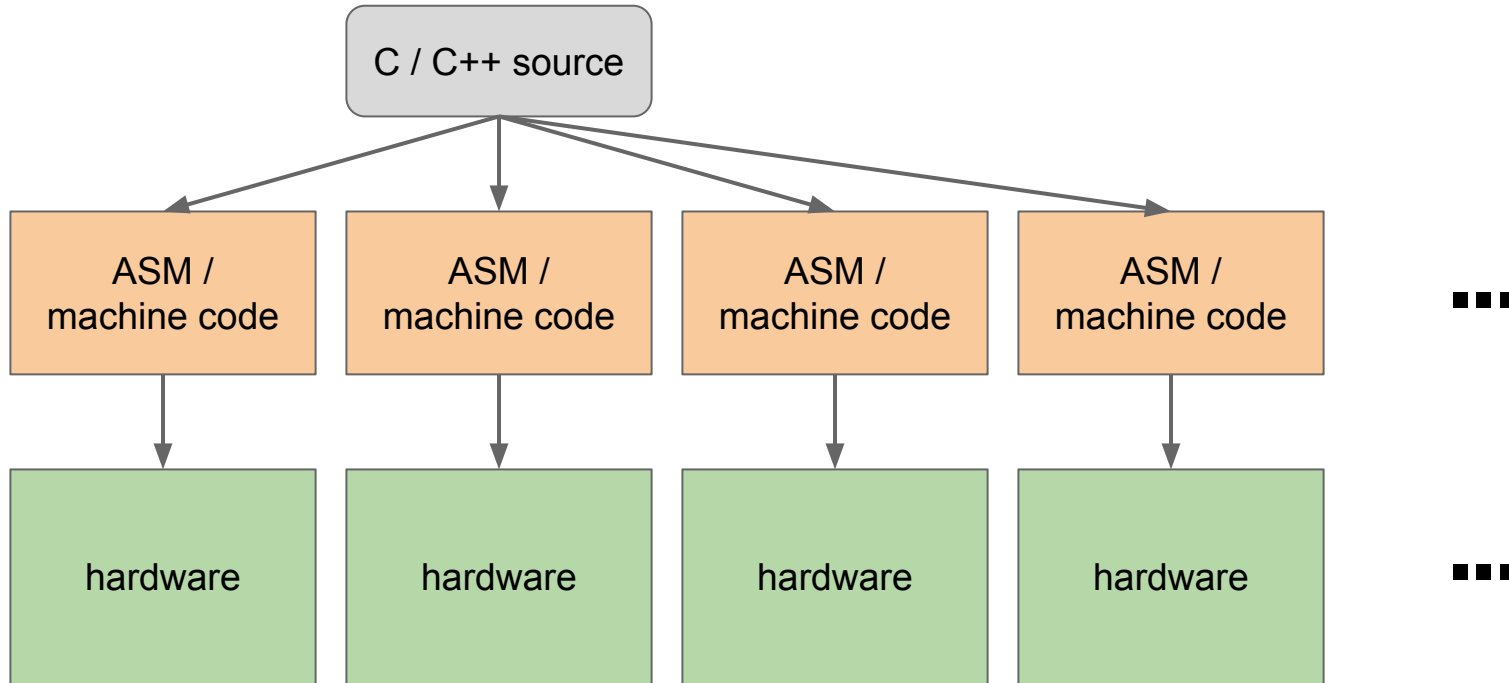
# Traditional (compiled) language vs. Java



# Traditional (compiled) language vs. Java

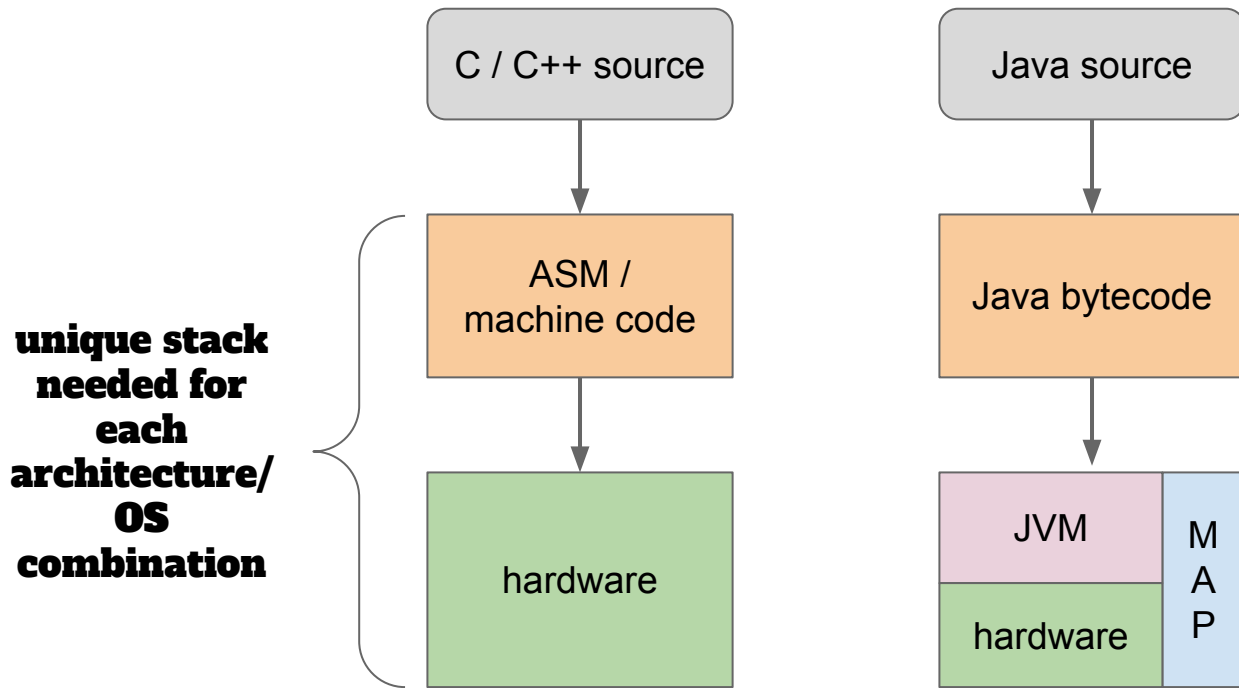


# Traditional (compiled) language vs. Java

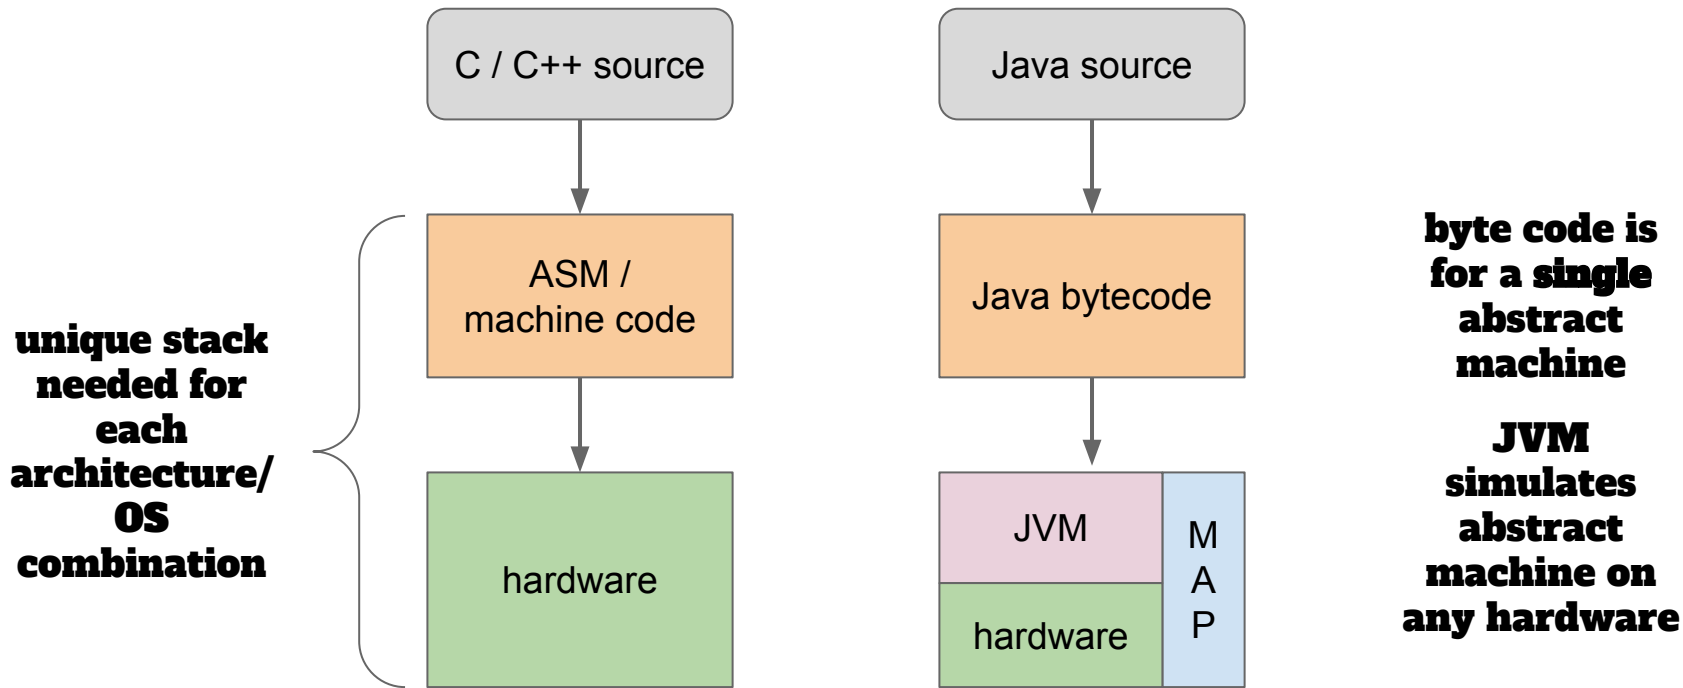





# Traditional (compiled) language vs. Java



# Traditional (compiled) language vs. Java

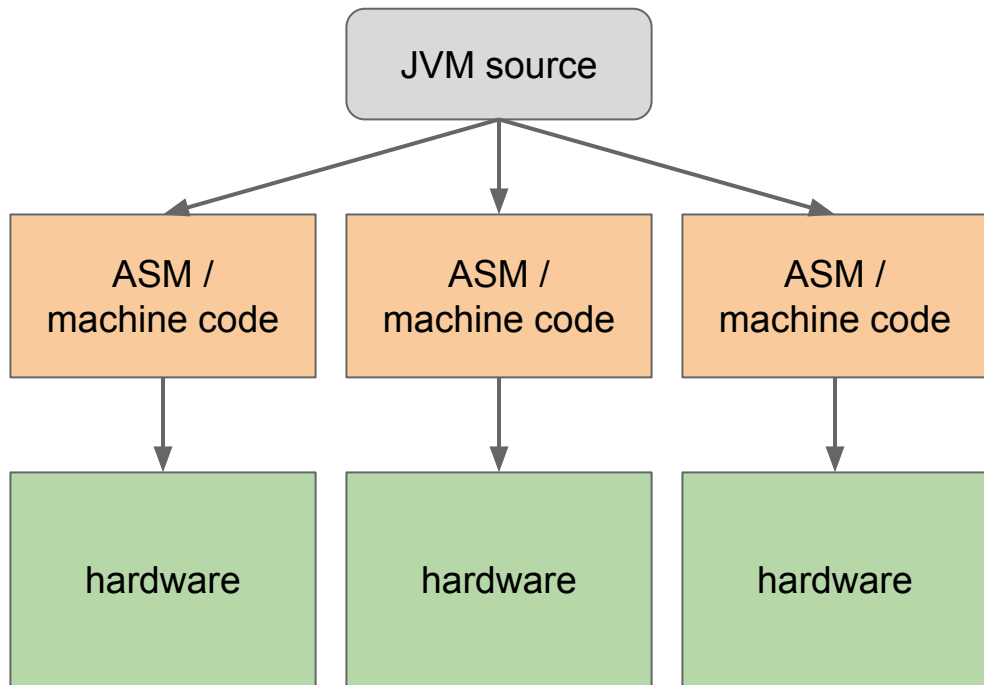


# JVM




- JVM is a program
- Written in C
- Compiled for each platform

# JVM




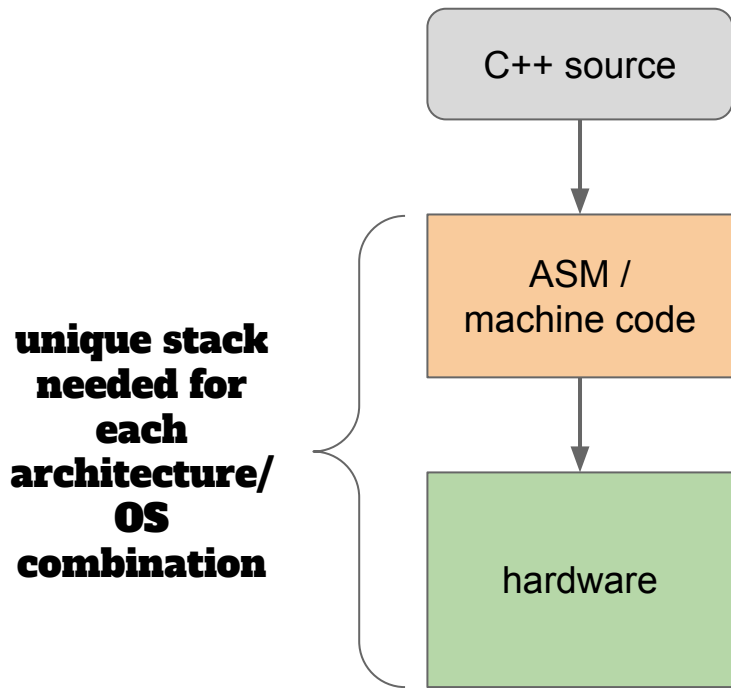
- JVM is a program
- Written in C
- Compiled for each platform
- Isn't that the same problem?

# JVM



- JVM is a program
- Written in C
- Compiled for each platform
- Isn't that the same problem?
  - Not really
  - One JVM source
  - C compilers exist for target hardware and OS.
  - We just have to build the JVM once for each
  - Then, **all** java programs can be run on that platform.

# Traditional (compiled) language vs. Java




**byte code is  
for a **single**  
abstract  
machine**

**JVM  
simulates  
abstract  
machine on  
any hardware  
MAP specific  
to arch/OS**

# Stack machine

- No registers
- Operands
  - Source: pop from stack
  - Destination: push onto stack
- Example
  - $C = A + B$

```
push A  
push B  
add  
pop C
```



# Stack machine

- No registers
- Operands
  - Source: pop from stack
  - Destination: push onto stack
- Example
  - $C = A + B$

push A  
push B  
add  
pop C

A



# Stack machine

- No registers
- Operands
  - Source: pop from stack
  - Destination: push onto stack
- Example
  - $C = A + B$

push A  
push B  
add  
pop C

A
B

# Stack machine


- No registers
- Operands
  - Source: pop from stack
  - Destination: push onto stack
- Example
  - $C = A + B$

```
push A  
push B  
add  
pop C
```

A
B


# Stack machine

- No registers
- Operands
  - Source: pop from stack
  - Destination: push onto stack
- Example
  - $C = A + B$



# Stack machine


- No registers
- Operands
  - Source: pop from stack
  - Destination: push onto stack
- Example
  - $C = A + B$



# Stack machine

- No registers
- Operands
  - Source: pop from stack
  - Destination: push onto stack
- Example
  - $C = A + B$

```
push A  
push B  
add  
pop C
```



# Why a stack machine?

- Compact object code
  - Operands are often implicit (from the stack)
  - So, “add”  
rather than “add x y”
- Less complex
  - No complex addressing modes
  - So, “inc”  
rather than “inc [si+bx-7]”

# Why a stack machine?

- Compact object code
  - Operands are often implicit (from the stack)
  - So, “add”  
rather than “add x y”
- Less complex
  - No complex addressing modes
  - So, “inc”  
rather than “inc [si+bx-7]”
- The primary reason: portability!

# Architectures differ wildly in register design

- x86
  - 8 16-bit registers
  - 8 8-bit registers
  - 4 segment registers
  - Instruction-specific registers
    - Implied operands
    - Restrictions
- ARM
  - 16 32-bit registers
  - 13 are general purpose
    - Technically you can do what you want with any of them
  - 3 have specific purpose

## Suppose JVM used registers?

- Which one architecture?
- What if more registers in JVM than architecture? Fewer?
- How to handle register sizes?
- How to manage special registers?

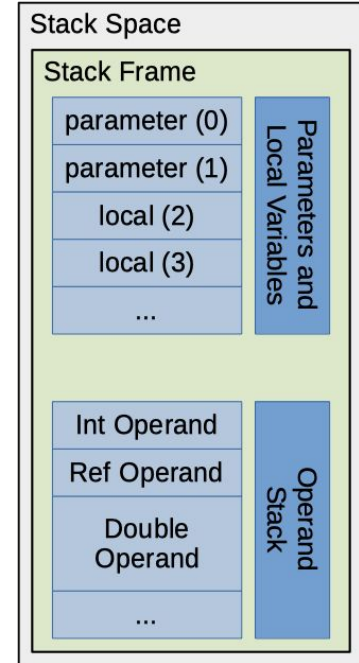

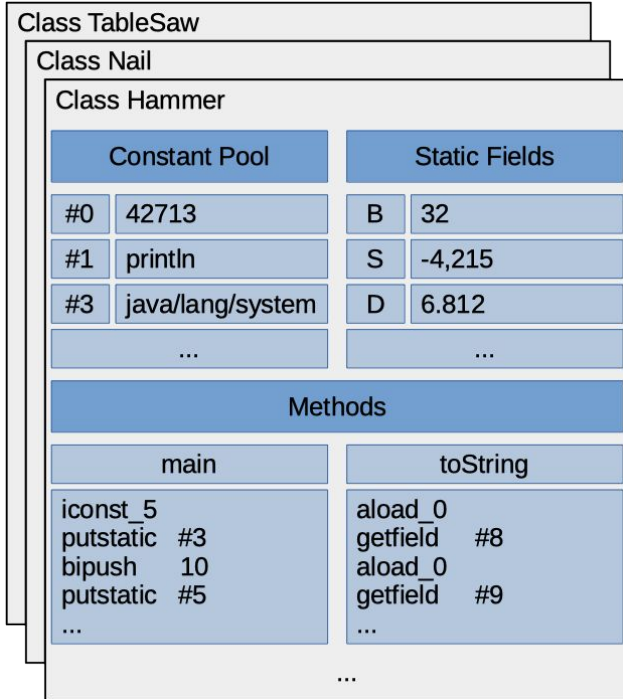


# Cost

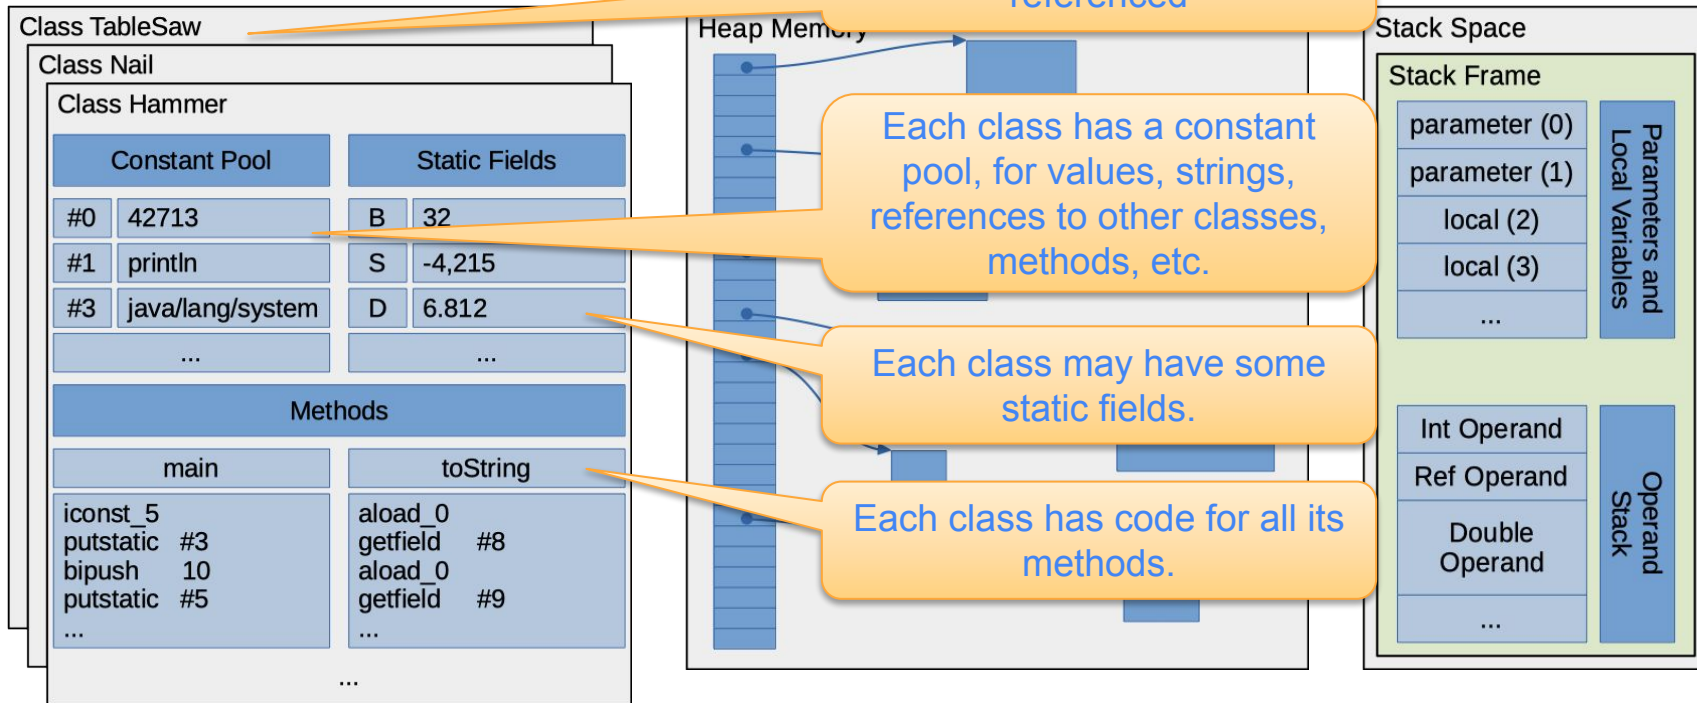
- All operands are in memory
  - Memory is slow
  - Registers are fast
- Registers don't require caching
  - Register can be temporary storage
  - Fast to access
  - No extra instructions to load/store

A = A+B	mov ax, [B]	push A
C = C+B	add [A], ax	push B
D = D+B	add [C], ax	add
	add [D], ax	pop A
		push C
		push B
		add
		pop C
		push D
		push B
		add
		pop D

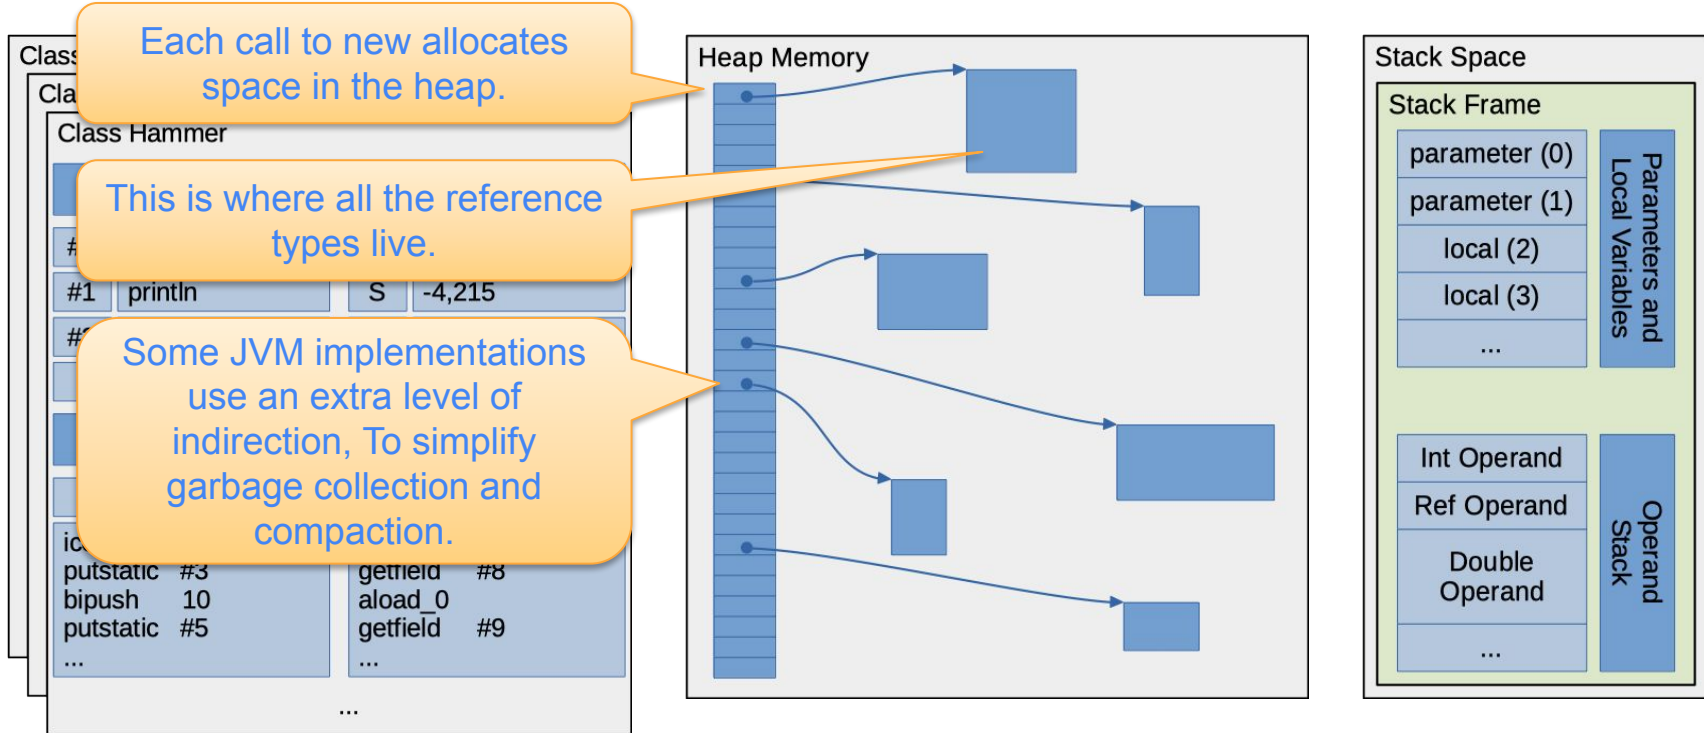
# JVM is an architecture



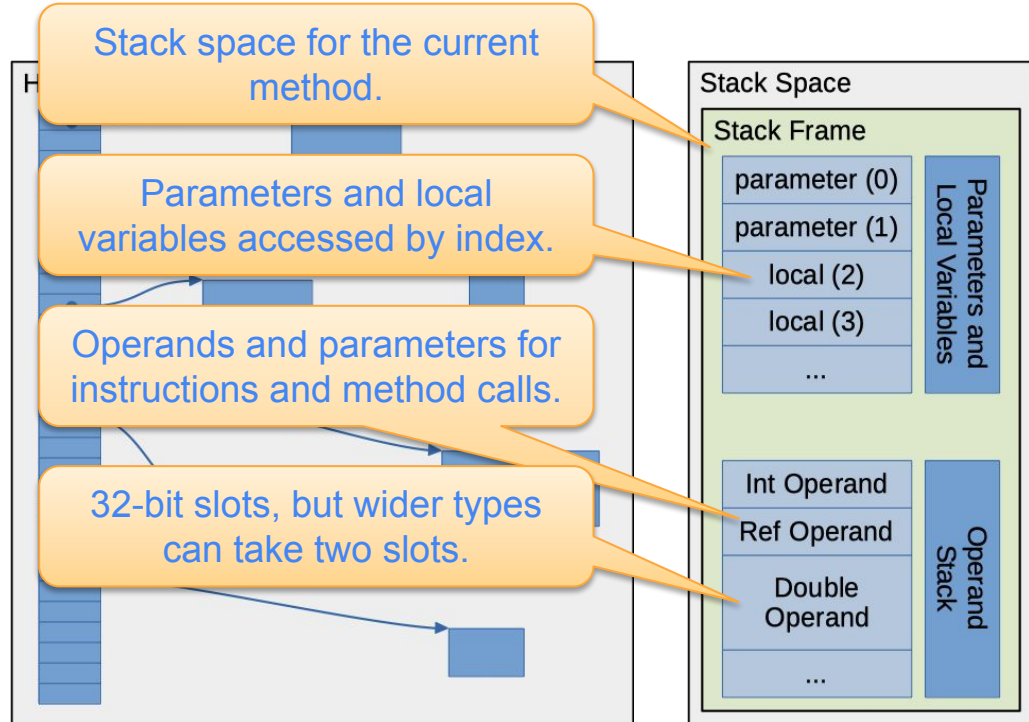
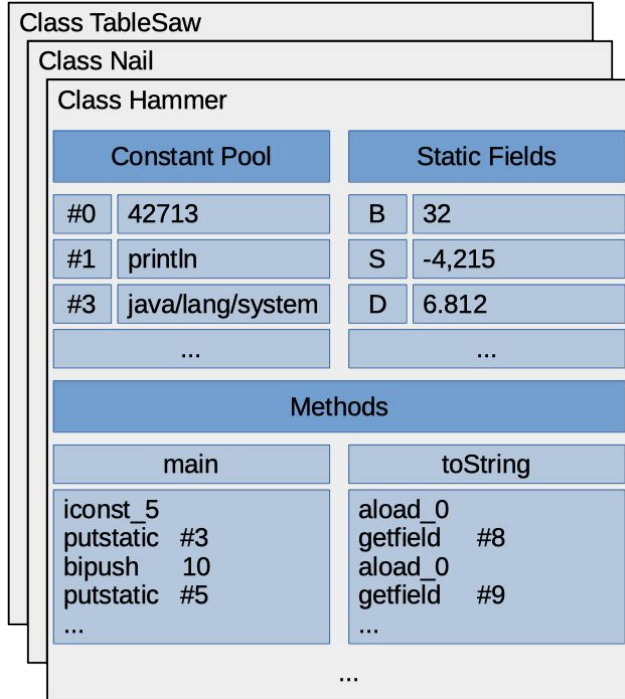
# JVM is an architecture



# JVM is an architecture




# JVM is an architecture



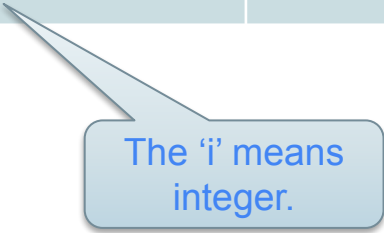
# Instructions

- One-byte opcode, up to 256 instructions
  - 202 defined so far, room for some extension
  - We'll learn more than 80 of them ... really fast.
- Variable length — most are 1 or 2 bytes
- Shortest instructions to access first 4 variables
- Short instructions to access up to 256 variables, constants, fields
- Longer, general instruction can access 65,536 constants, fields



# Meet Some Bytecode

Instruction	Operand (size)	Description
iload_0		Push integer value of local variable 0 onto stack.
iload_1		Push integer value of local variable 1 onto stack.
iload_2		Push integer value of local variable 2 onto stack.
iload_3		Push integer value of local variable 3 onto stack.
iload	<i>idx</i> (8)	Push value of local variable <i>idx</i> onto stack.



The 'i' means integer.

# Loading Literal Values

Instruction	Operand (size)	Description
iconst_m1		Push -1 onto the stack
iconst_0		Push 0 onto the stack.
iconst_1		Push 1 onto the stack.
iconst_2		Push 2 onto the stack.
...		...
iconst_5		Push 5 onto the stack.
bipush	<i>val</i> (8)	Push byte <i>val</i> as an int
bspush	<i>val</i> (16)	Push short <i>val</i> as an int
ldc	<i>idx</i> (8)	Load the value at index <i>idx</i> from the constant pool.



# Instructions for Integer Math

Instruction	Operand (size)	Description
iadd		Pop integers B then A, push $A + B$
isub		Pop integers B then A, push $A - B$
imul		Pop integers B then A, push $A * B$
idiv		Pop integers B then A, push $A / B$
irem		Pop integers B then A, push $A \% B$

iand,ior, ixor		Pop integers B then A, push $A \& B$ , $A   B$ or $A \wedge B$
ishl, ishr		Pop integers B then A, push $A \ll B$ or $A \gg B$

# Saving Results

Instruction	Operand (size)	Description
istore_0		Pop integer value and store in local variable 0.
istore_1		Pop integer value and store in local variable 1.
istore_2		Pop integer value and store in local variable 2.
istore_3		Pop integer value and store in local variable 3.
ilstore	<i>idx</i>	Pop integer value and store in local variable at <i>idx</i> .

# Operands for Different Types

Integer	Long	Float	Double
iload_0 .. 3	lload_0 .. 3	fload_0 .. 3	dload_0 .. 3
iload	lload	fload	dload
iconst_m1 .. 5	lconst_0 .. 1	fconst_0 .. 2	dconst_0 .. 1
istore_0 .. 3	lstore_0 .. 3	fstore_0 .. 3	dstore_0 .. 3
istore	lstore	fstore	dstore
iadd, isub ...	<i>ladd, lsub ...</i>	fadd, fsub ...	dadd, dsub ...

- Where are the instructions for bool, byte or short?
  - Local operands and local variables are really ints.
  - Array elements and fields can be these smaller types
  - .. but they're converted to int when we operate on them.

# More Instructions

- We're skipping lots of instructions
  - Stack manipulation
  - Type conversion
  - Field and array access
  - Conditionals, branching and jumping
  - Calling methods and returning

<b><u>Instruction type</u></b>	<b><u>%</u></b>
<b>Local-variable push</b>	<b>34.5</b>
<b>Local-variable pop</b>	<b>7.0</b>
<b>Memory push</b>	<b>20.2</b>
<b>Memory pop</b>	<b>4.0</b>
<b>Constant push</b>	<b>6.8</b>
<b>Compute</b>	<b>9.2</b>
<b>Branches</b>	<b>7.9</b>
<b>Calls/returns</b>	<b>7.3</b>
<b>Misc stack ops</b>	<b>2.1</b>
<b>New objects</b>	<b>0.4</b>
<b>All others</b>	<b>0.6</b>

## **Survey of Bytecodes**


**Note : More pushes  
than pops**

**C = A + B**

**A is local 0**

**B is local 1**

**C is local 2**



**frame**


**C = A + B**

**A is local 0**

**B is local 1**

**C is local 2**

iload	0



**frame**


**C = A + B**

**A is local 0**

**B is local 1**

**C is local 2**

iload	0
iload	1




**frame**



**C = A + B**  
**A is local 0**  
**B is local 1**  
**C is local 2**

iload	0
iload	1
iadd	



**frame**


**C = A + B**

**A is local 0**

**B is local 1**

**C is local 2**


iload	0
iload	1
iadd	
istore	2



**frame**

**C = A + B**  
**A is local 0**  
**B is local 1**  
**C is local 2**

iload	0
iload	1
iadd	
istore	2



**frame**

iload	0	iload	1	iadd	istore	2
-------	---	-------	---	------	--------	---


**C = A + B**  
**A is local**  
**0**  
**B is local**  
**1**  
**C is local**  
**2**

iload_0	
iload_1	
iadd	
istore_2	

A shorter  
encoding.

Less important with  
Just-In-Time  
Compilation

iload_0	iload_1	iadd	istore_2
---------	---------	------	----------




**frame**

```
static void f( int a ) {  
    int b = ?;  
    int c = ?;  
    int d = ?;  
    int e;  
  
    e = a + b * c;  
}
```


Parameters are  
the first local  
variables.

this for an  
instance  
method.



```
static void f( int a ) {  
    int b = ?;  
    int c = ?;  
    int d = ?;  
    int e;  
  
    e = a + b * c;  
}
```

Then the local  
variables.




```

static void f( int a ) {
  int b = ?;
  int c = ?;
  int d = ?;
  int e;

  e = a + b * c;
}

```

iload_0	



**frame**


```

static void f( int a ) {
  int b = ?;
  int c = ?;
  int d = ?;
  int e;

  e = a + b * c;
}

```

iload_0	
iload_1	
iload_2	



**frame**




```

static void f( int a ) {
  int b = ?;
  int c = ?;
  int d = ?;
  int e;

  e = a + b * c;
}

```

iload_0	
iload_1	
iload_2	
imul	



**frame**


```

static void f( int a ) {
  int b = ?;
  int c = ?;
  int d = ?;
  int e;

  e = a + b * c;
}

```

iload_0	
iload_1	
iload_2	
imul	
iadd	



**frame**


```

static void f( int a ) {
  int b = ?;
  int c = ?;
  int d = ?;
  int e;

  e = a + b * c;
}

```

iload_0	
iload_1	
iload_2	
imul	
iadd	
istore	4



**frame**

# Parameter passing

```
int calc(int a, int b)
```


How is this done?

In ASM/C pass parameters on stack

Same in Java


# Parameter passing

**locals**




**fixed, static size**

**operands**





# Parameter passing

**locals**




**operands**




# Parameter passing


**locals**



**operands**




**push params  
on stack**




# Parameter passing



**locals**



**operands**




**call  
procedure**






# Parameter passing


**locals**



**operands**




**call  
procedure**




# Parameter passing


**locals**



**operands**



**callee's  
local  
frame**



**caller's  
operand  
stack**



# Meet javap

- The java development kit includes **javap**
- ... a class file disassembler.
- It can show you what's in a class file.
- With the `-c` option, it can show bytecode for each method:

```
$ javap -c Test.class
public class Test {
    public static void main(java.lang.String[]);
    Code:
        0: ldc           #2      // String Hello
        2: astore_1
        3: ldc           #3      // String world.
        5: astore_2
        6: getstatic     #4      // Field ...
```

That's java bytecode  
right there.

# Performance

- Java does not run *natively*
  - All instructions are *interpreted* by JVM
- Executing an instruction
  - Fetch: read next instruction
  - Decode: determine what to do
  - Execute: do it
- For assembly, that's one machine instruction.
- How long in JVM?
  - A single byte code instruction becomes
  - 10 to 100s of machine instructions (e.g. new instruction)

# Executing Java bytecode

## (1) Software interpreter

- Interpret each instruction
- Perform required operations
  - Change the state of the JVM
- Short loop to execute JVM instructions one after another
  - Repeatedly fetch & interpret the next instruction
  - Same process over and over
  - ... but operands and JVM state will vary from instruction to instruction.

## Speedup techniques

- Cache code
  - Saves disk/file IO

# Executing Java bytecode

## (2) Compile to native code

- Instead of emitting bytecode
  - Emit native ASM/machine code
- Rarely done
- Why
  - Defeats portability
  - Lots of other aspects of the JVM (e.g., garbage collection) that still need to be provided.
  - Difficult to optimize

## Speedup techniques

- Native code for individual JVM instruction
  - fetch/decode/execute done by the processor
  - ... not done via software.
- Still has overhead
  - Stack operations (rather than directly using registers)

# Executing Java bytecode

## (3) JIT (just in time) compiler

- Read class files for each new class
  - Including bytecode for each method.
- Individual methods can be compiled to sequence of native instructions.
  - If they are expensive or executed frequently.

## Speedup techniques

- Native methods can run faster
- Not all methods need to be compiled to native instructions
- Can even support different implementations of the same method.
- For example:
  - `double(x) ⇒ add ax,ax`
  - `double(2) ⇒ mov ax,4`

# Executing Java bytecode

## (4) JVM hardware chip

- Build hardware to run bytecode natively
- Architecture is designed specifically for bytecode
  - Faster than general-purpose machine

## Speedup techniques

- Very fast
  - For JVM
  - Bytecode is inefficient compared to traditional
  - Eg, no registers
- Disadvantages
  - Special-purpose hardware can be expensive.