



Advanced architecture






42 Years of Microprocessor Trend Data



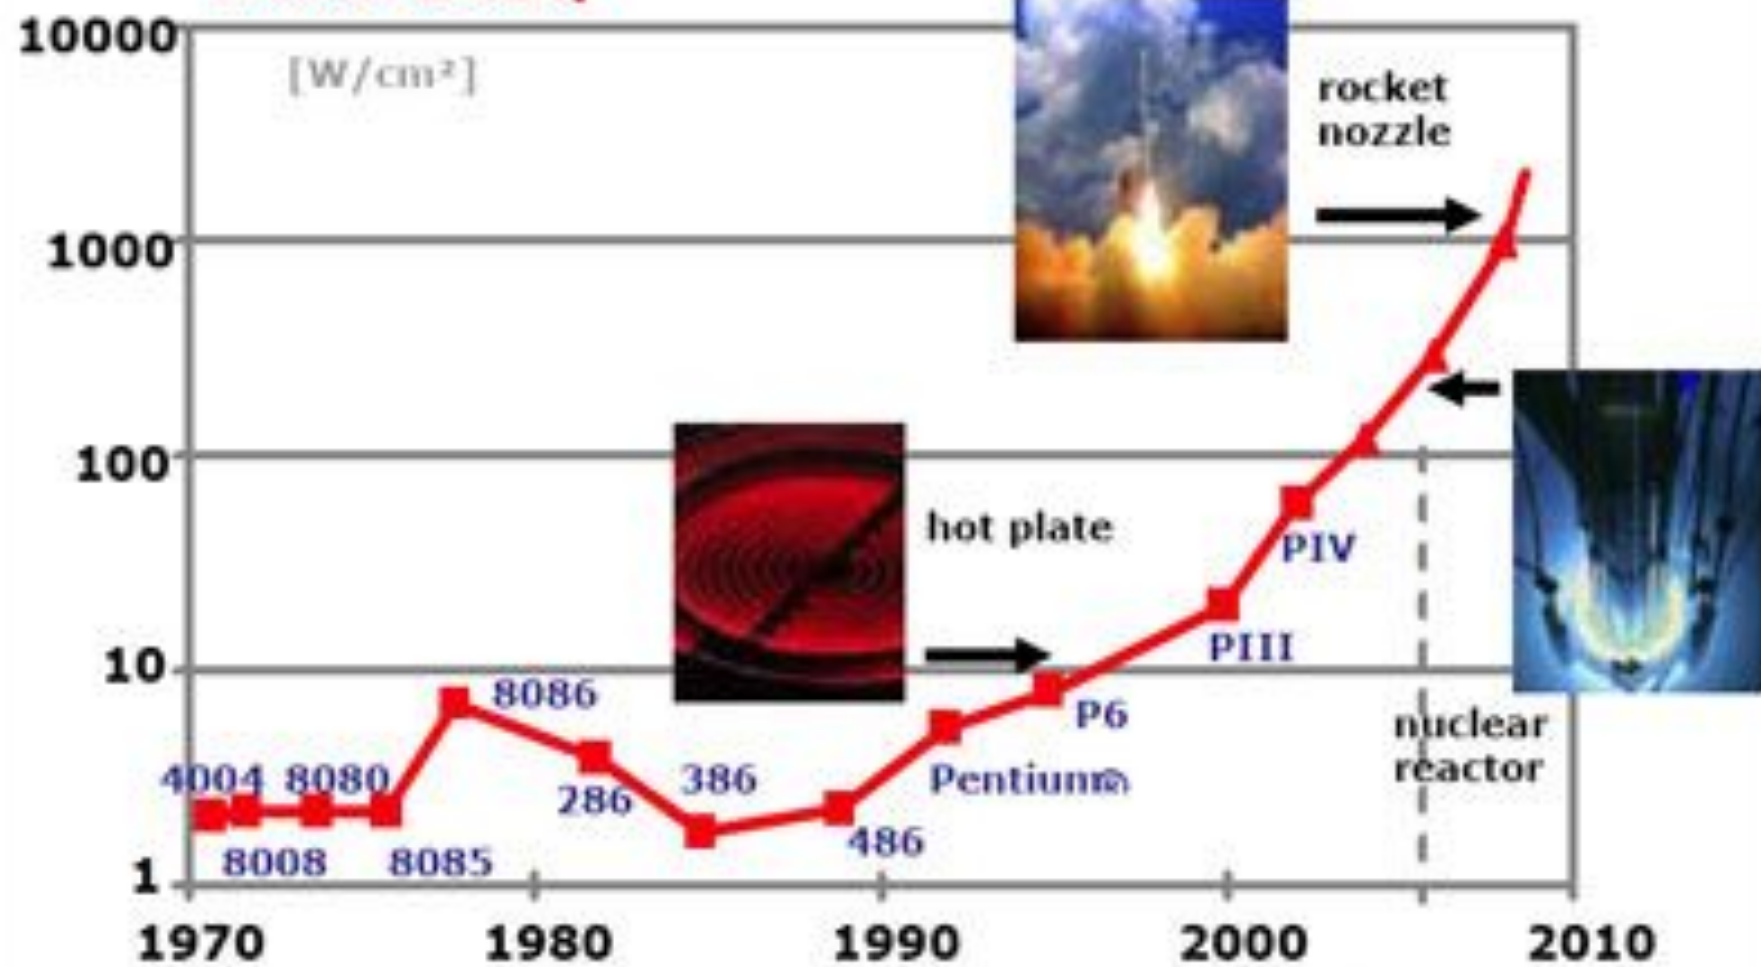
Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

50 years of microprocessor size (from wikipedia)



AMD Ryzen 7 3700X (64-bit, SIMD, caches, I/O die)	5,990,000,000 ^{[124][d]}	2019	AMD	7 & 12 nm (TSMC)
HiSilicon Kirin 990 4G	8,000,000,000 ^[125]	2019	Huawei	7 nm
Apple A13 (iPhone 11 Pro)	8,500,000,000 ^{[126][127]}	2019	Apple	7 nm
AMD Ryzen 9 3900X (64-bit, SIMD, caches, I/O die)	9,890,000,000 ^{[1][2]}	2019	AMD	7 & 12 nm (TSMC)
HiSilicon Kirin 990 5G	10,300,000,000 ^[128]	2019	Huawei	7 nm
AWS Graviton2 (64-bit, 64-core ARM-based, SIMD, caches) ^{[129][130]}	30,000,000,000	2019	Amazon	7 nm
AMD Epyc Rome (64-bit, SIMD, caches)	39,540,000,000 ^{[1][2]}	2019	AMD	7 & 12 nm (TSMC)
Apple M1	16,000,000,000 ^[131]	2020	Apple	5 nm

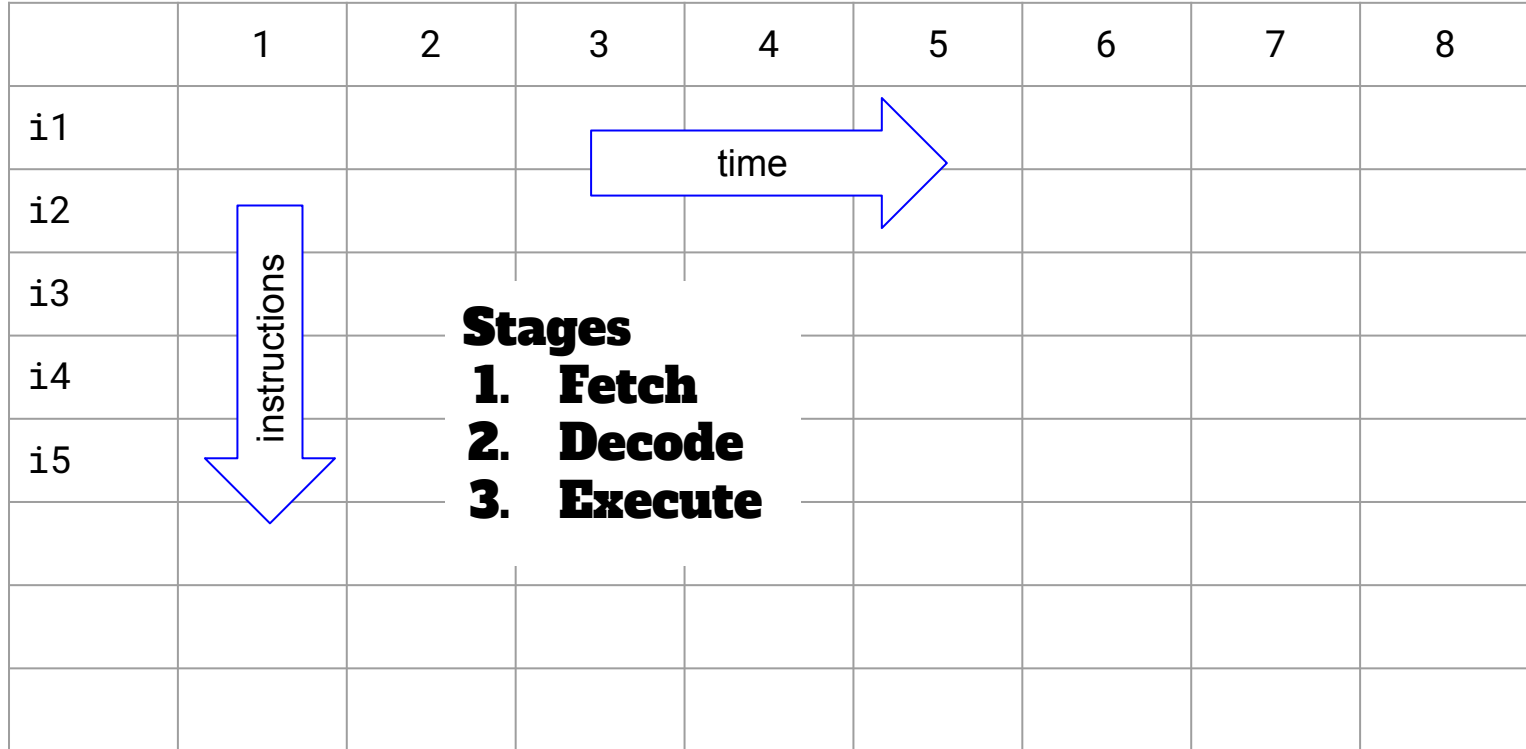
Power density



So many transistors!

- What are we doing with so many transistors?
- Lots of clever things
 - Lots of cache
 - Pipelining
 - Branch prediction
 - Delayed branching
 - Superscalar execution
 - Out-of-order execution
 - Multiple cores
 - Hyper-threading

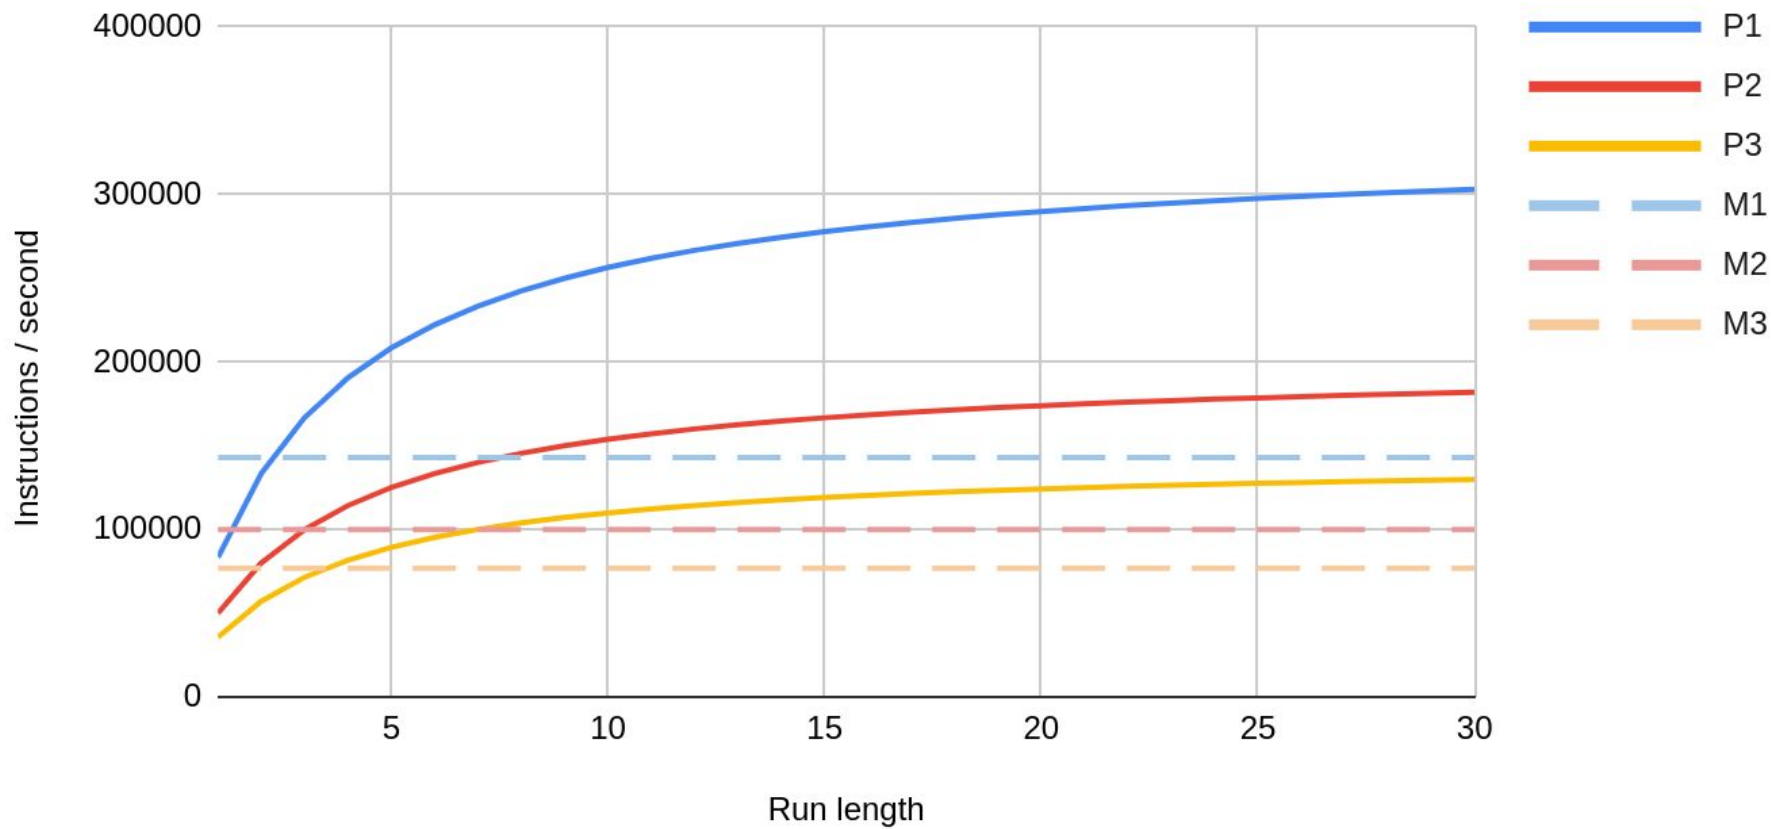
Pipelining example



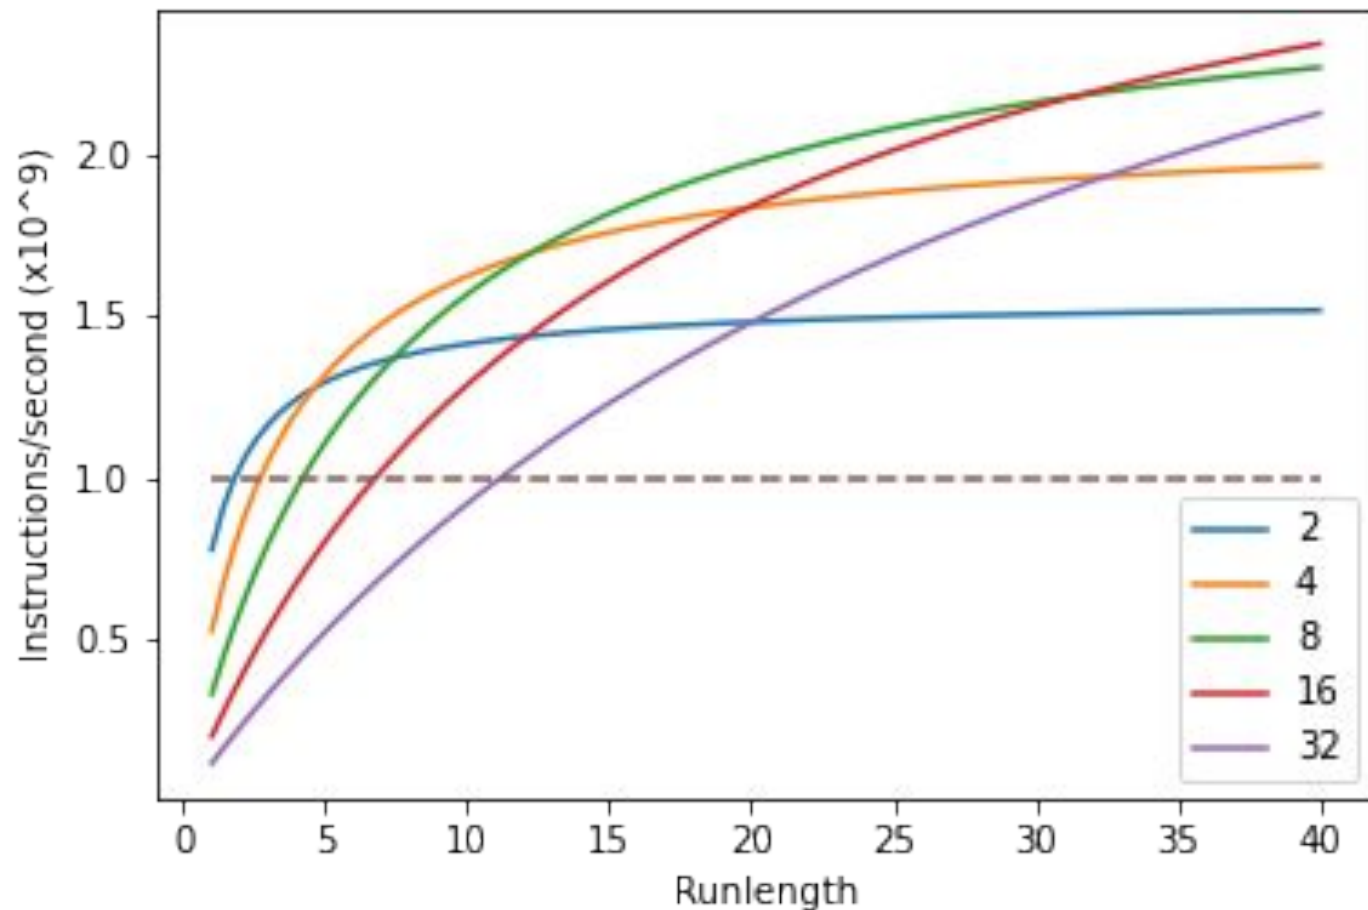
Pipelining example

	1	2	3	4	5	6	7	8
i1	F	D	E					
i2		F	D	E				
i3			F	D	E			
i4				F	D	E		
i5					F	D	E	
i6						F	D	E

Pipelined Performance



Pipeline length effect on performance



$$IPS = \frac{L \cdot f(S)}{L + S - 1} \times 10^9$$

$$t(I) = f(S) \cdot C$$

Branch prediction

- Want to have long run lengths
 - What to do on branch?
 - Four choices

Stop

- **Drains pipeline**
- **Poor performance**

Guess

- **Right: no penalty**
- **Wrong:**
 - **Same as stopping**
 - **Well, some cost**

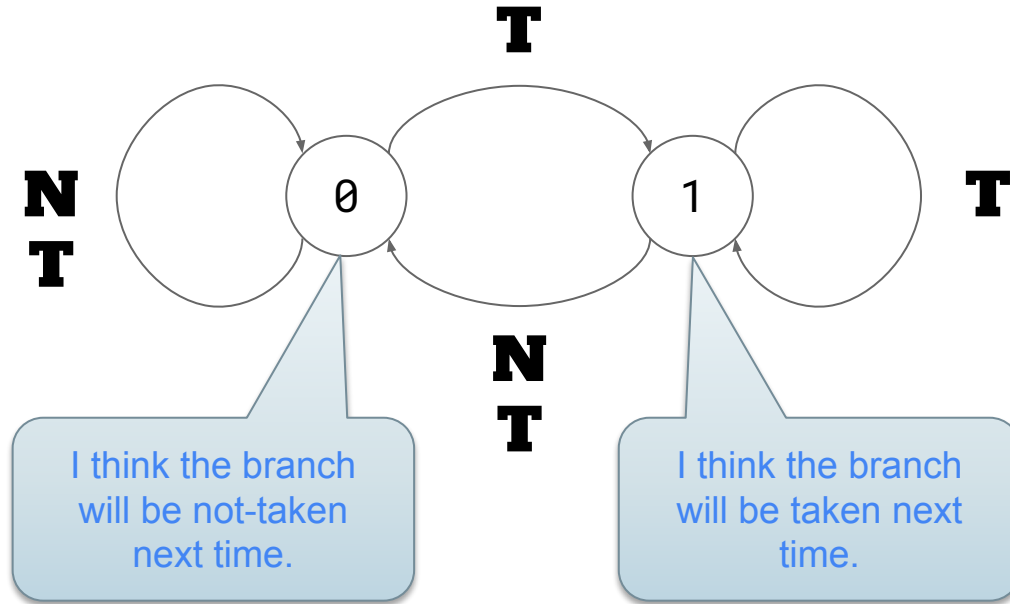
Predict

- **Same as guess**
- **Improved accuracy**


Fetch both

- **Best performance**
- **Highest cost**
- **Limited**

1-bit branch predictor



2-bit branch predictor



Delayed branching

```
repit:  cmp ax,0
        je  iszero    ;ax=0
        add ax,1
        jmp repit
```

...

```
iszero:                ;ax=1
```

- Instruction following conditional jump is **always** executed

Delay slot

```
rep: cmp [si], 0
      jge skip
      ??????????????
      mov [si], 0
skip:
      inc si
      loop rep
```


```
rep: cmp [si], 0
      jge skip
      inc si
      mov [si-1], 0
skip:
      loop rep
```

**Must move an
instruction into the
*delay slot***

Delayed branching

	1	2	3	4	5	6	7	8
mov	F	D	E					
add		F	D	E				
jo			F	D	E			
inc				F	D			
mov					F			

Delayed branching



	1	2	3	4	5	6	7	8
add	F	D	E					
jo		F	D	E				
mov			F	D	E			
inc				F	D			
mov					F			

ILP — instruction-level parallelism

Three statements (instructions)

$e = a + b$

$f = c + d$

$g = e * f$



independent

Can execute first two in either order or in parallel

$e = a + b$ $f = c + d$

$g = e * f$

- **Superscalar execution : finish two or more instructions in the same cycle.**

ILP is 3/2 for these instructions.

Superscalar

	1	2	3	4	5	6	7	8
i1	F	D	E	S				
i2	F	D	E	S				
i3		F	D	E	S			
i4		F	D	E	S			
i5			F	D	E	S		
i6			F	D	E	S		
i7				F	D	E	S	
i8				F	D	E	S	

Multiple functional units

- For two-way superscalar must have 2 of everything
 - Two fetch units
 - Two decode units
 - Two execute units, including all parts of the ALU
 - ...

Multiple functional units

- For two-way superscalar must have 2 of everything
 - Two fetch units
 - Two decode units
 - Two execute units, including all parts of the ALU
 - ...
- Can we really build a processor with duplicate hardware?
Yes. Yes we can.
- Even the the 8086 has multiple functional units
 - Consider `add ax, [si+dx]`
 - How many adders?

Superscalar execution

- Some pairs of instructions won't occur in a typical program

```
jmp    left  
jmp    right
```

```
cmp     ax, 25  
cmp     ax, 30
```

```
inc     bx  
dec     bx
```



- Some pairs of instructions may have interdependencies
 - Preventing them from being run concurrently.

```
add     ax, [var1]  
add     bx, ax
```

- Where can we get lots of code with no interdependency between instructions?

Running Multiple Threads



- Instructions in one program (or thread) are typically independent of those in another thread.



- This provides this is an opportunity for parallelism.
- Support execution of multiple instruction streams at the same time.

Multiple Cores

- Modern processors have multiple CPU cores.
 - Like having multiple processors on the same package.



- We could accomplish this with multiple CPU packages.
- But that would require fabricating more than one CPU
- Multiple motherboard slots ...

Multi-threading

- Threads may not always keep their CPU busy
- ... because of memory stalls, unused parts of the ALU, etc.



Average utilization = 60%



Average utilization = 100%

Multi-threading



Highly idealized, ie, false

Overly idealized, ie, false

There is a cost to switch between threads

Hyperthreading

- Intel's name for *simultaneous multithreading*
- First, what is multithreading?
 - OS executes processes or threads
 - Multiple concurrent threads
 - OS *swaps* between threads
- Why?

Content switch

- Thread context
 - Process info (PID, owner, ...)
 - Page table, TLB
 - Stack
 - Microprocessor state (registers, ...)
- Context switch
 - Save current thread's context
 - Restore next thread's context
- Secondary cost
 - Flush/refresh cache (especially the TLB)
 - Page buffers

Multi-threading



Highly idealized, ie, false



Overly idealized, ie, false





Hyperthreading


- Normally, the OS switches between running threads
 - When one thread has to wait for something (e.g., I/O).
 - But, some waits are too short for this to be practical
- Idea : let the CPU trade-off between a small number of running threads
- The CPU can pretend it has more cores than it really does
 - The OS can schedule these *virtual cores* as if they were real CPU cores
 - Internally, the processor can switch between them when one can't make progress.
- Intel calls this *hyper-threading*.

Hyperthreading


- The CPU can maintain a limited number of contexts for executing threads
 - Copies of the CPU registers
 - Cache contents
- OS sees each of these as if it was a CPU core.
- Internally, they share ALU and other hardware.







Cinebench 11.5 multi-threaded



Handbrake Video Encoding Test HT

Custom PC Benchmark

bit-tech.net



Intel Core i7-8700K



Intel Core i7-7700K



References

- https://www.researchgate.net/figure/Evolution-of-multi-core-processors-1_fig1_281534326
- <https://www.extremetech.com/extreme/203490-moores-law-is-dead-long-live-moores-law>
- <https://www.karlsruhp.net/2018/02/42-years-of-microprocessor-trend-data/>
- <http://www.nanowerk.com/spotlight/spotid=1762.php>
- <http://www.pctechguide.com/pentium-cpus/hyper-threading-technology>
- <https://www.makeuseof.com/tag/hyperthreading-technology-explained/>
- https://en.wikipedia.org/wiki/Transistor_count